

Lecture 26

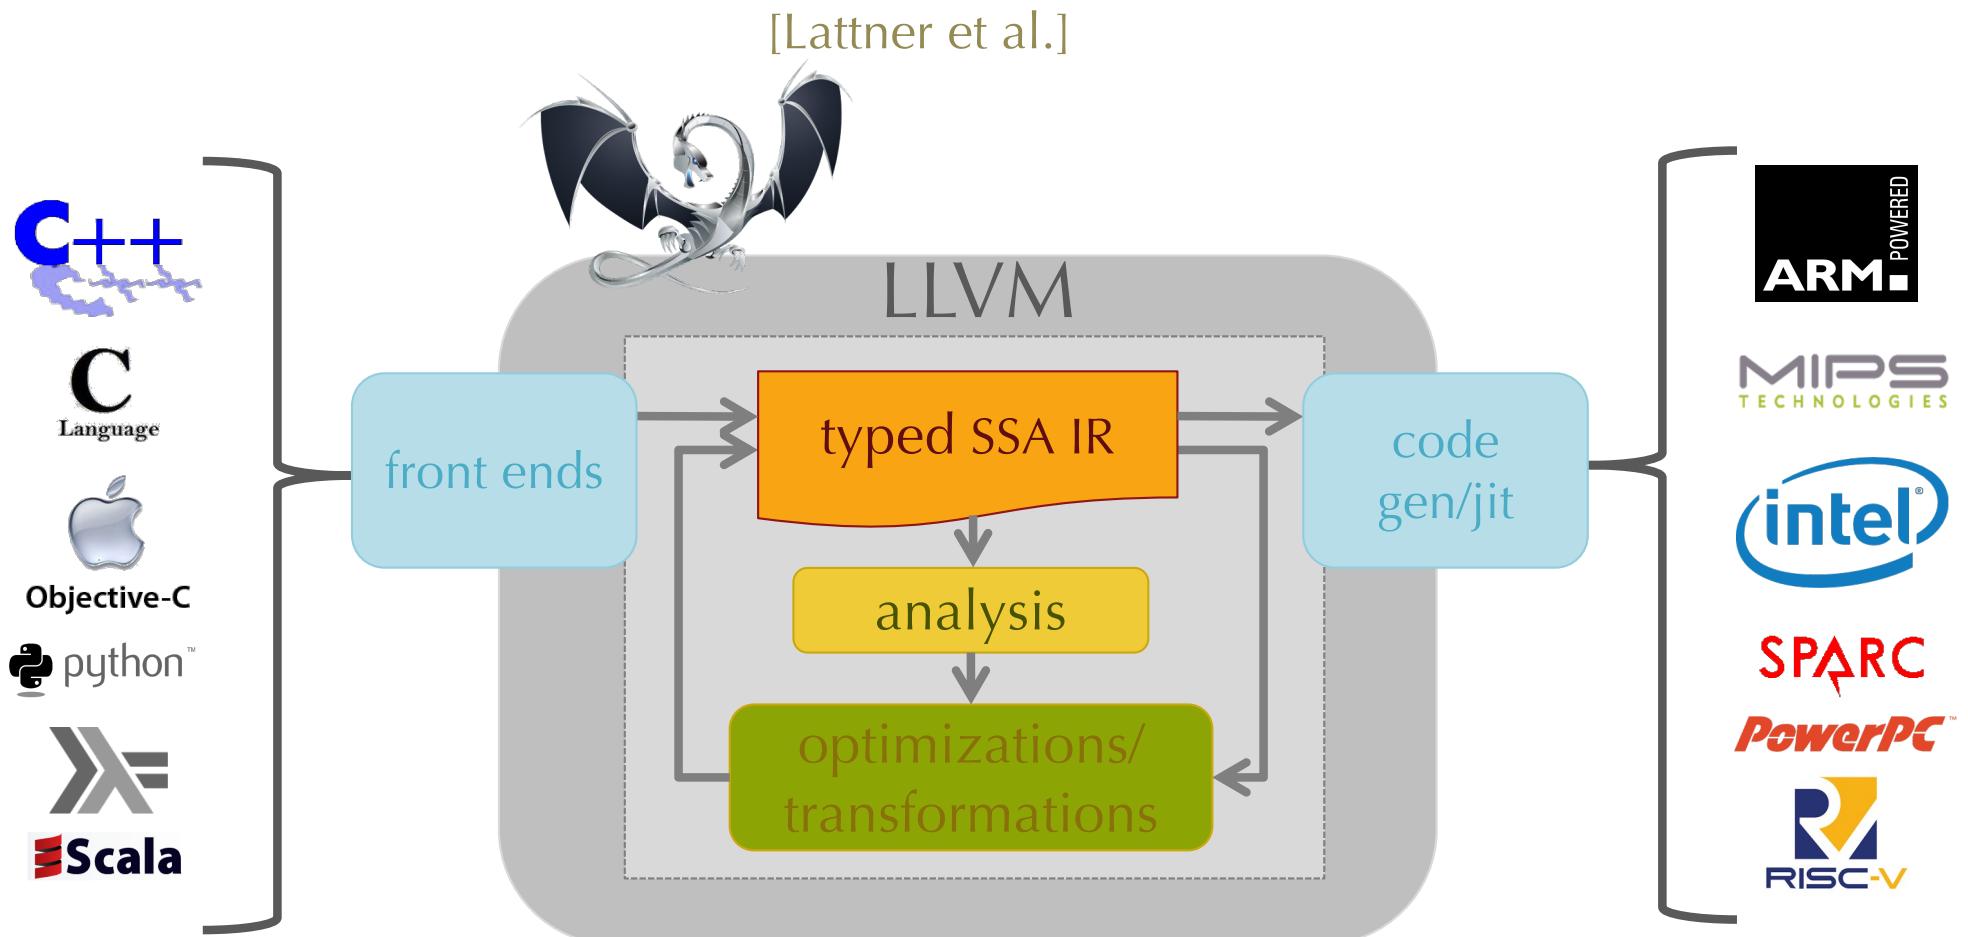
CIS 341: COMPILERS

Announcements

- Final Exam:
 - LRSM AUD
 - Monday, May 2nd noon - 2:00pm
- Current Plan / My Preference: *In Person*
 - Unless University policy prohibits in person exams, this is the default
 - If you have serious concerns about taking the exam in person, I will make accommodations

COMPILER VERIFICATION

LLVM Compiler Infrastructure



Other LLVM IR Features

- C-style data values
 - ints, structs, arrays, pointers, vectors
- Type system
 - used for layout/alignment/padding
- Relaxed-memory concurrency primitives
- Intrinsics
 - extend the language malloc, bitvectors, etc.
- Transformations & Optimizations



Make targeting LLVM IR
easy and attractive for
developers!

But... it's complex



One Example: **undef**

The **undef** "value" represents an arbitrary, but indeterminate bit pattern for any type.

Used for:

- uninitialized registers
- reads from volatile memory
- results of some underspecified operations

What is the value of `%y` after running the following?

```
%x = or i8 undef, 1
%y = xor i8 %x, %x
```

One plausible answer: 0
Not LLVM's semantics!

(LLVM is more liberal to permit more aggressive optimizations)

Partially defined values are interpreted *nondeterministically* as sets of possible values:

```
%x = or i8 undef, 1  
%y = xor i8 %x, %x
```

$$[\![\text{i8 undef}]\!] = \{0, \dots, 255\}$$

$$[\![\text{i8 1}]\!] = \{1\}$$

$$\begin{aligned} [\![\%\text{x}]\!] &= \{a \text{ or } b \mid a \in [\![\text{i8 undef}]\!], b \in [\![1]\!]\} \\ &= \{1, 3, 5, \dots, 255\} \end{aligned}$$

$$\begin{aligned} [\![\%\text{y}]\!] &= \{a \text{ xor } b \mid a \in [\![\%\text{x}]\!], b \in [\![\%\text{x}]\!]\} \\ &= \{0, 2, 4, \dots, 254\} \end{aligned}$$

Interactions with Optimizations

Consider:

```
%y = mul i8 %x, 2
```

versus:

```
[%x] = [i8 undef]  
= {0,1,2,3,4,5,...,255}  
[%y] = {a mul 2 | a ∈ [%x]}  
= {0,2,4,...,254}
```

```
%y = add i8 %x, %x
```

```
[%x] = [i8 undef]  
= {0,1,2,3,4,5,...,255}  
[%y] = {a + b | a ∈ [%x],  
b ∈ [%x]}  
= {0,1,2,3,4,...,255}
```



Interactions with Optimizations

Consider:

```
%y = mul i8 %x, 2
```

versus:

```
%y = add i8 %x, %x
```

Upshot: if **%x** is **undef**, we
can't optimize **mul** to **add**
(or vice versa)!

What's the problem?

Bug List: (12 of 435) First Last Prev Next Show last search results

Bug 33165 - Simplify* cannot distribute instructions for simplification due to undef

Status: REOPENED

Reported: 2017-05-25 02:12 PDT by Nuno Lopes

Davide Italiano 2017-05-25 08:55:40 PDT

[Comment 6](#)

W:

cc:

To:

no (unless we want to give up on some undef transformations, and special case selection
but I'm afraid others might be affected too)

By: John Regehr 2017-05-25 09:09:24 PDT

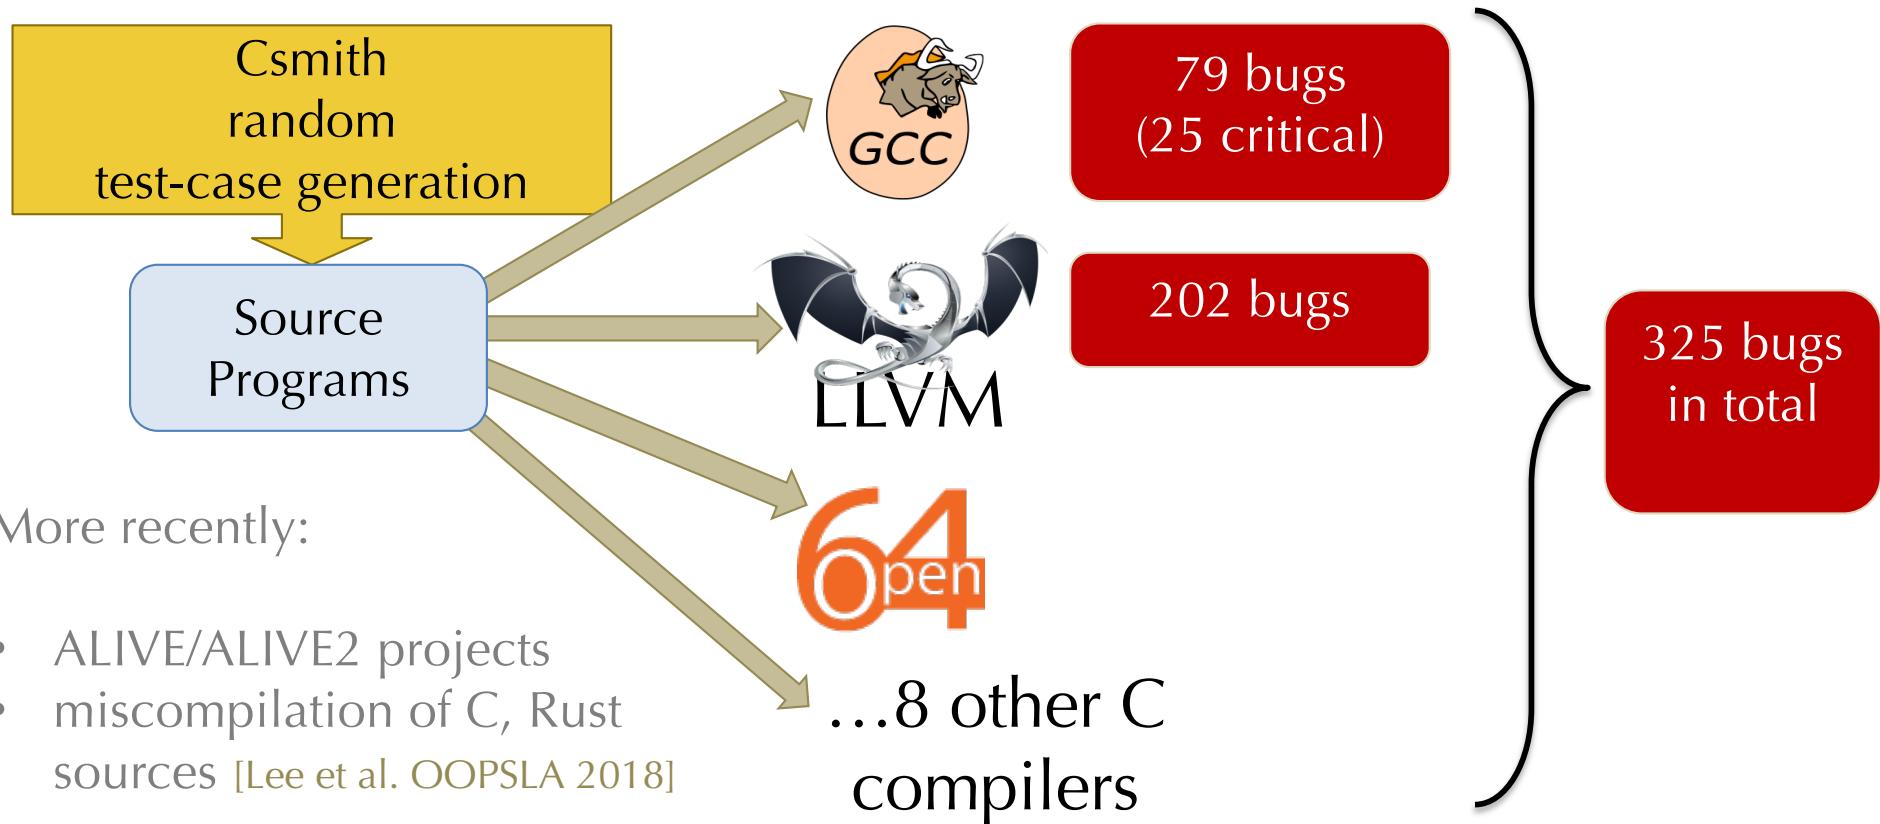
[Comment 7](#)

Yes, this is one of those test cases. There are so many optimization failures
Nuno has been automatically filtering out classes of mistranslation that are
to be hard to fix but I guess he decided to take a closer look at some of the

Soon I'll be able to include branches/phis in these test cases, but only for
branches due to a limitation in Alive.

Compiler Bugs

[Regehr's group: Yang et al. PLDI 2011]

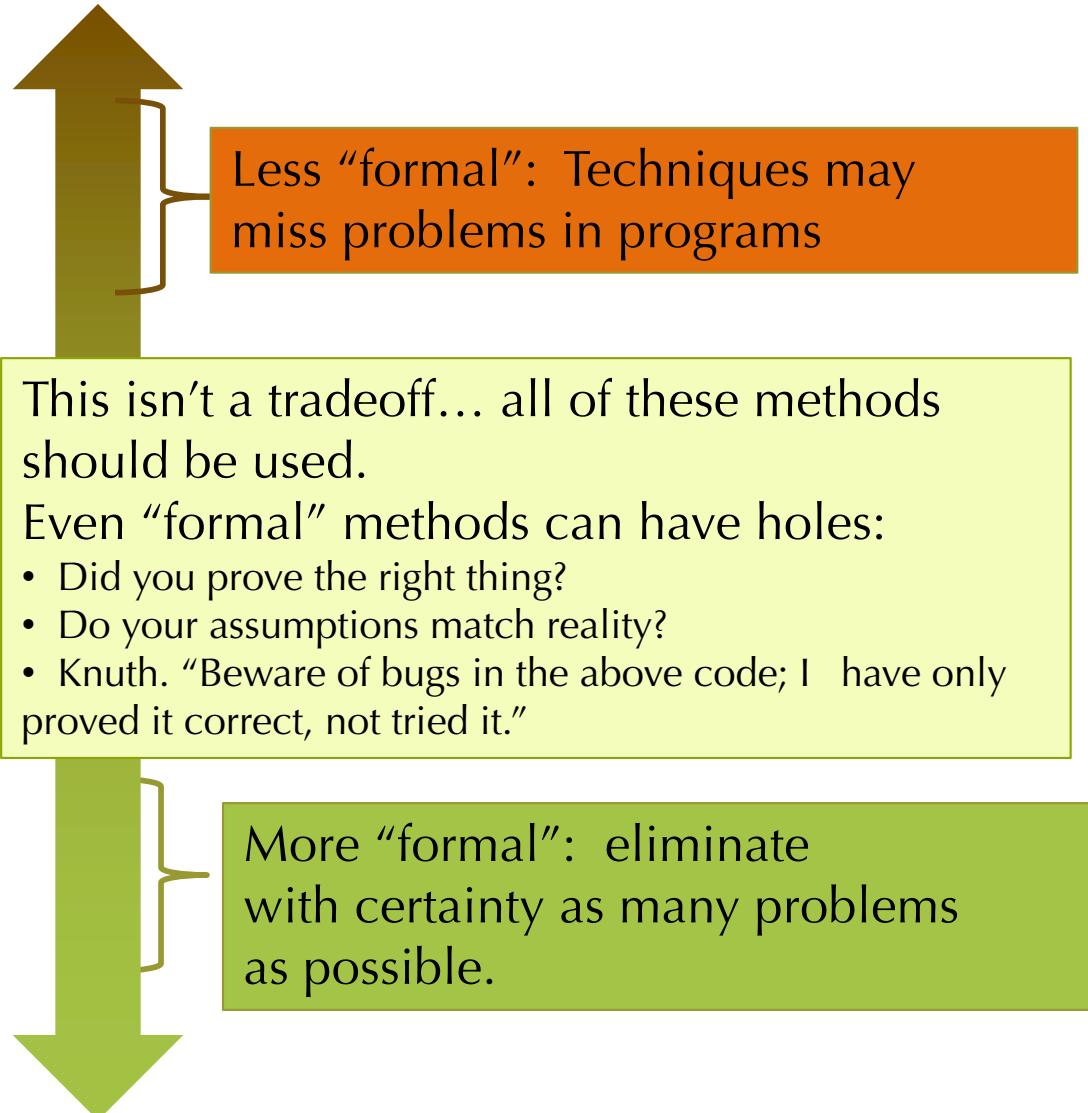


LLVM is hard to trust
(especially for critical code)

What can we do about it?

Approaches to Software Reliability

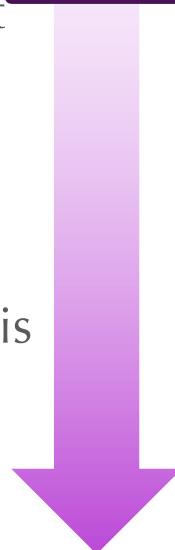
- Social
 - Code reviews
 - Extreme/Pair programming
- Methodological
 - Design patterns
 - Test-driven development
 - Version control
 - Bug tracking
- Technological
 - “lint” tools, static analysis
 - Fuzzers, random testing
- Mathematical
 - Sound programming languages tools
 - “Formal” verification



Goal: Verified Software Correctness

- Social
 - Code reviews
 - Extreme/Pair programming
- Methodological
 - Design patterns
 - Test-driven development
 - Version control
 - Bug tracking
- Technological
 - “lint” tools, static analysis
 - Fuzzers, random testing
- Mathematical
 - Sound programming languages tools
 - “Formal” verification

Q: How can we move the needle towards mathematical software correctness properties?



Taking advantage of advances in computer science:

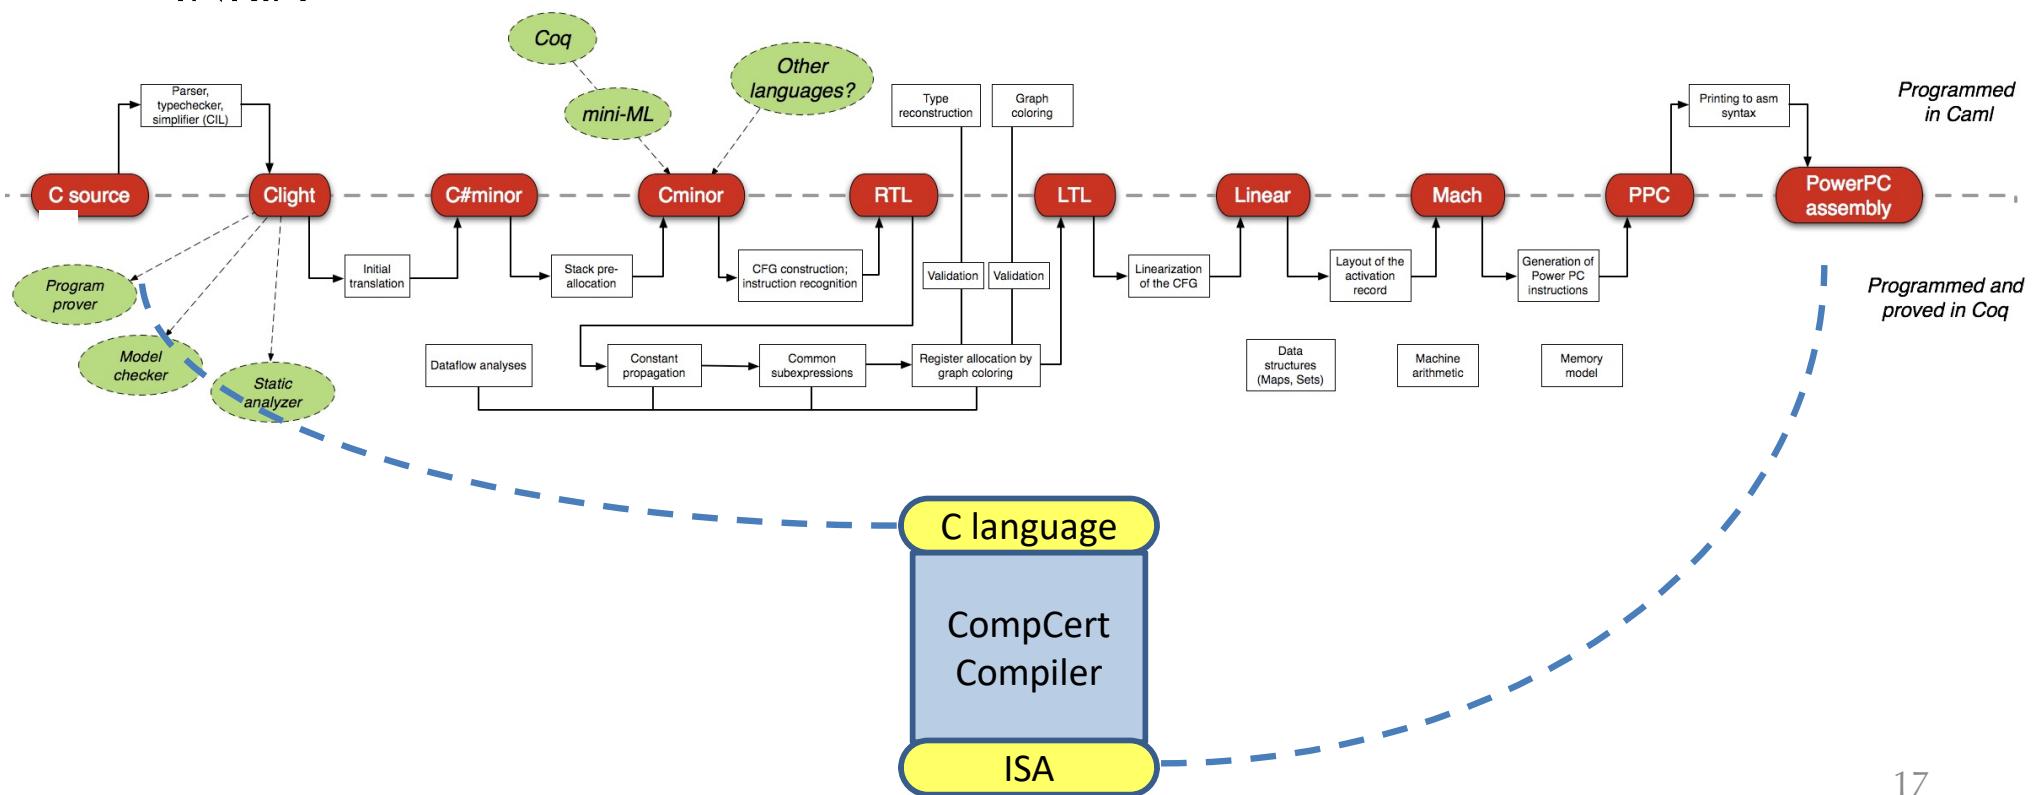
- Moore's law
- improved programming languages & theoretical understanding
- better tools:
interactive theorem provers

CompCert – A Verified C Compiler



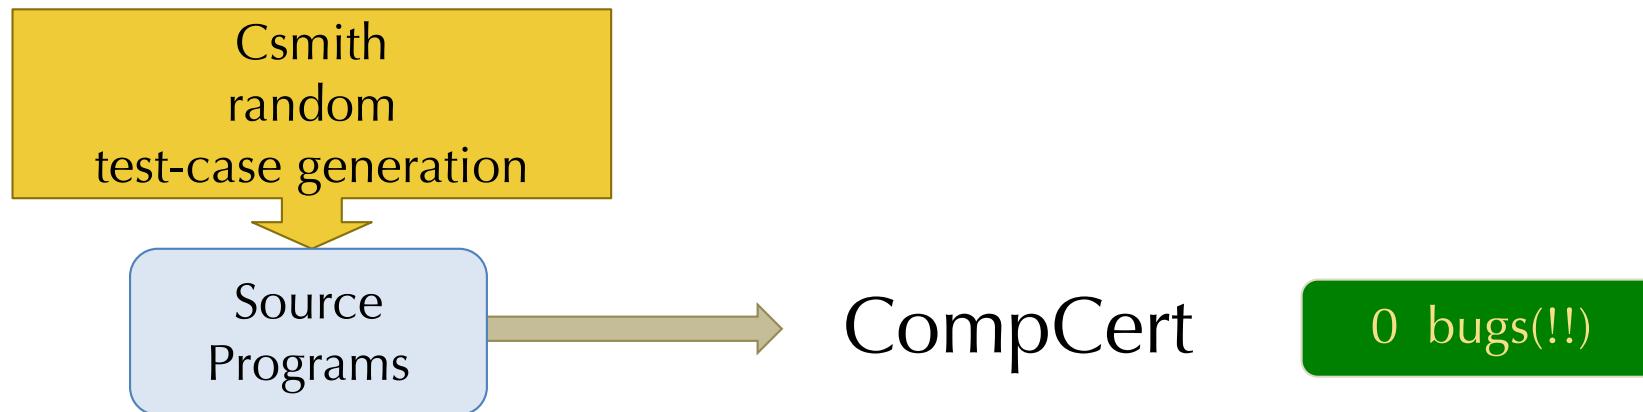
Xavier Leroy
INRIA

Optimizing C Compiler,
proved correct end-to-end
with machine-checked proof in Coq



Csmith on CompCert?

[Yang et al. PLDI 2011]



Verification Works!

"The striking thing about our CompCert results is that the middle-end bugs we found in all other compilers are absent. As of early 2011, the under-development version of CompCert is the only compiler we have tested ***for which Csmith cannot find wrong-code errors***. This is not for lack of trying: we have devoted about six CPU-years to the task. ***The apparent unbreakability of CompCert supports a strong argument that developing compiler optimizations within a proof framework, where safety checks are explicit and machine-checked, has tangible benefits for compiler users.***"

– Regehr et. al 2011

Our Approach: Formal Verification

Interactive theorem proving in Coq

- not model checking / SMT
- human-in-the-loop



Using Coq **is** functional programming
...but some of your programs *are* proofs

⇒ proof engineering



deepspec.org

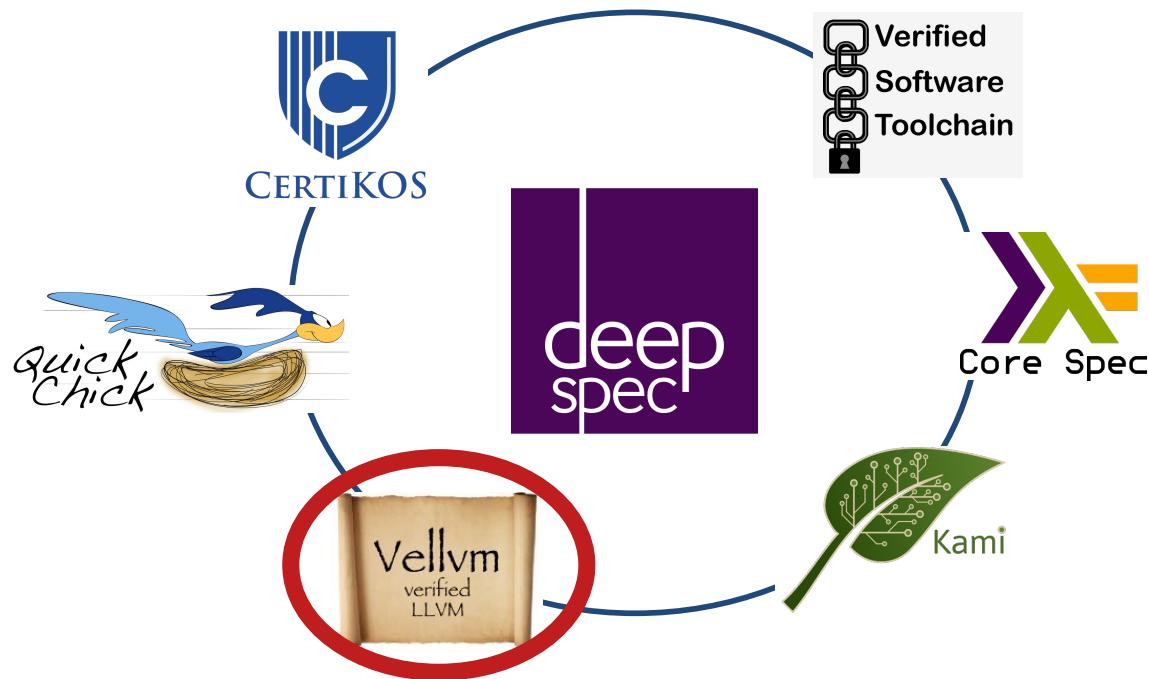
Deep Specifications

[deepspec.org]



- ***Rich*** – expressive description
- ***Formal*** – mathematical, machine-checked
- ***2-Sided*** – tested from both sides
- ***Live*** – connected to real, executable code

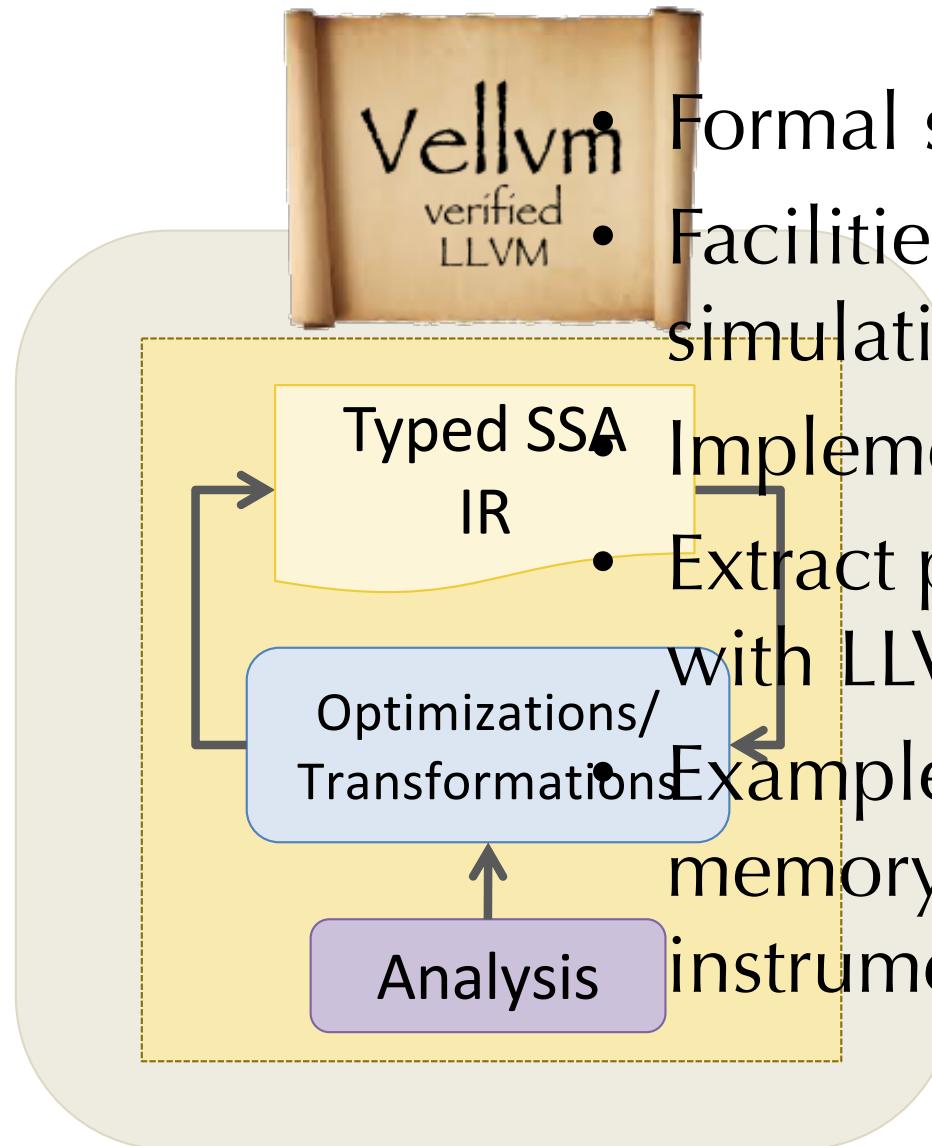
Goal: Advance the reliability, safety, security, and cost-effectiveness of software (and hardware).



deepspec.org

The Vellvm Project

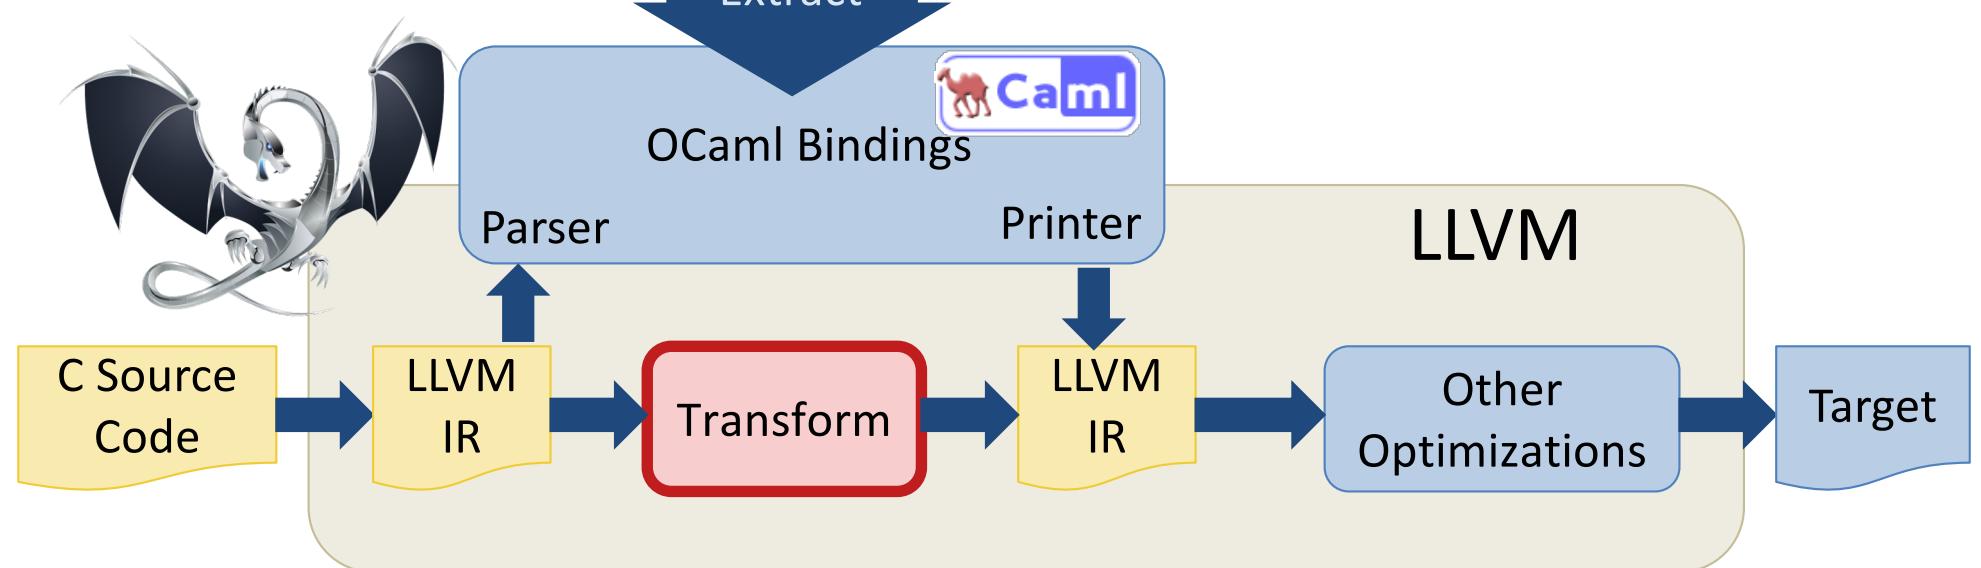
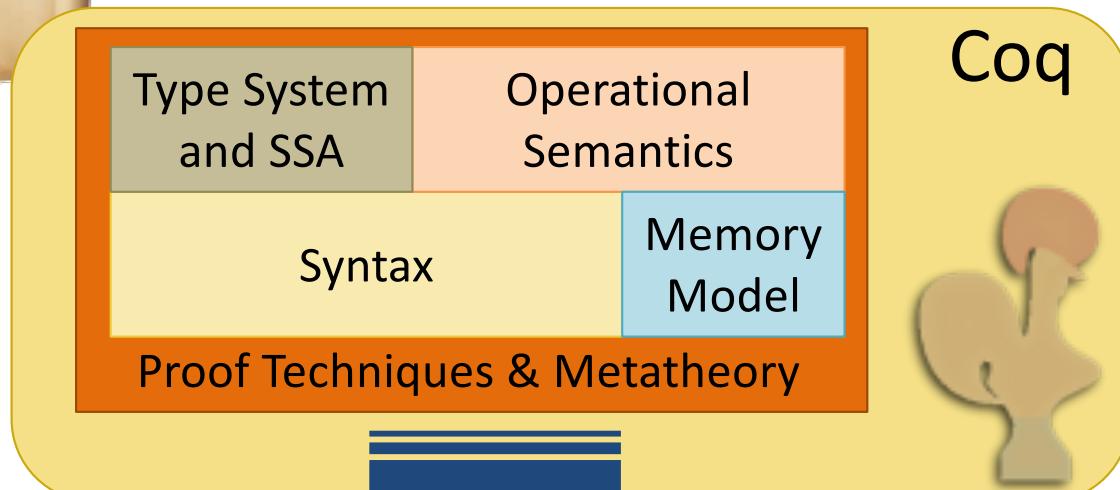
[Zhao et al. POPL 2012, CPP 2012, PLDI 2013, Zackowski, et al. ICFP2021]



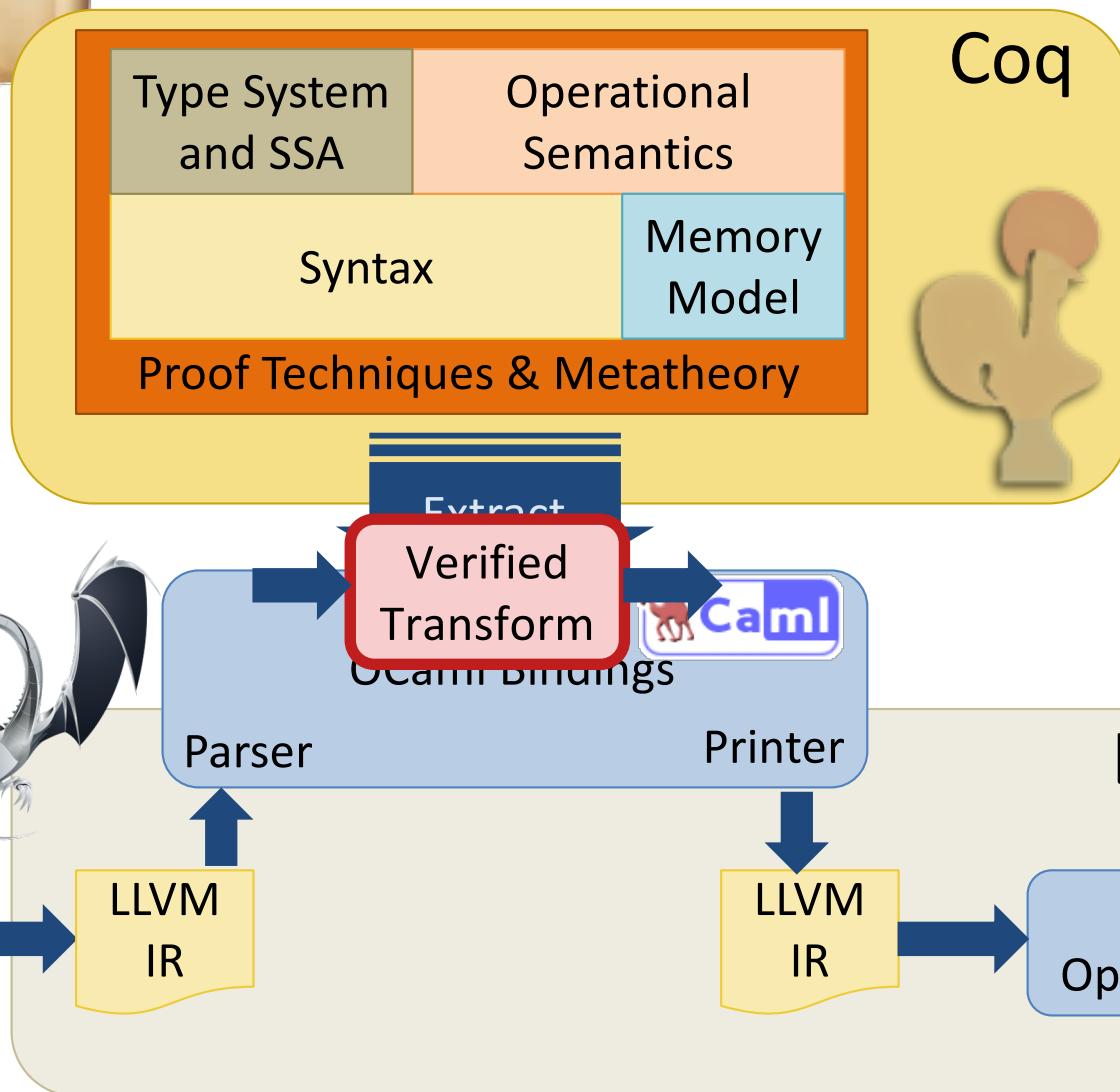
- Formal semantics
 - Facilities for creating simulation proofs
 - Implemented in Coq
 - Extract passes for use with LLVM compiler
- Example: verified memory safety instrumentation



Vellvm Framework



Vellvm Framework



Writing Interpreters in Coq



Galina (Coq's Language)

- rich, **dependent type** system
- **pure, total functional** language

How do we write the interpretation function?

```
Inductive exp : Set :=
| EXP_Ident (id:ident)
| EXP_Integer (x:int)
| EXP_Fraction (n:int) (* IN *)
| EXP_Rational (n:int) (* CO *)
| EXP_Negation (exp:exp)
| EXP_Plus (exp1:exp) (exp2:exp)
| EXP_Minus (exp1:exp) (exp2:exp)
| EXP_Mult (exp1:exp) (exp2:exp)
| EXP_Divide (exp1:exp) (exp2:exp)
| EXP_Pow (exp1:exp) (exp2:exp)
| EXP_Sqrt (exp:exp)
| EXP_Cond (exp1:exp) (exp2:exp) (exp3:exp)

Inductive instr : Set :=
| INSTR_Op (op:exp)
| INSTR_Call (fn:texp) (args:list texp)
| INSTR_Definition code := list (instr_id * instr).
| INSTR_Label label := string;
| INSTR_Goto label := string;
| INSTR_Exit := unit;
| INSTR_StackPush (texp) := texp;
| INSTR_StackPop (texp) := texp;
| INSTR_StackSwap (texp) := texp;
| INSTR_StackCopy (texp) := texp;
| INSTR_StackFill (texp) := texp;
| INSTR_StackClear (texp) := texp;
| INSTR_StackAlloc (texp) := texp;
| INSTR_StackFree (texp) := texp;
| INSTR_StackPushLabel (label) := label;
| INSTR_StackPopLabel (label) := label;
| INSTR_StackSwapLabel (label) := label;
| INSTR_StackCopyLabel (label) := label;
| INSTR_StackFillLabel (label) := label;
| INSTR_StackClearLabel (label) := label;
| INSTR_StackAllocLabel (label) := label;
| INSTR_StackFreeLabel (label) := label;

Record block : Set :=
mk_block
{
  blk_id      : block_id;
  blk_phis   : list (local_id * phi);
  blk_code    : code;
  blk_term   : instr_id * terminator;
}.
```

Datatypes for Abstract Syntax

LLMV Memory Model (simplified)

(* IO interactions for the LLVM IR *)

Inductive IO : Type -> Type :=

| Alloca : $\forall (t:\text{dtyp}), (\text{IO dvalue})$
| Load : $\forall (t:\text{dtyp}) (a:\text{dvalue}), (\text{IO dvalue})$
| Store : $\forall (a:\text{dvalue}) (v:\text{dvalue}), (\text{IO unit})$
| GEP : $\forall (t:\text{dtyp}) (v:\text{dvalue}) (vs:\text{list dvalue}), (\text{IO dvalue})$
| ItoP : $\forall (i:\text{dvalue}), (\text{IO dvalue})$
| PtoI : $\forall (a:\text{dvalue}), (\text{IO dvalue})$
| Call : $\forall (f:\text{string}) (\text{args:list dvalue}), (\text{IO dvalue})$
.

Describes the interface
for "observations" of
LLVM IR programs.

output values of
the Call event

type of the result
provided by the
environment

LLVM Interpreter in Coq

```
Definition step (s:state) : LLVMTrace result
let '(g, pc, e, k) := s in
do cmd ← trywith ("CFG has no instruction at " ++ string_of_pc pc)
match cmd with
| Term (TERM_Ret (t, op)) =>
'dv ← eval_exp (Some (eval_typ t)) op;
match k with
| [] => halt dv
| (KRet e' id p') :: k' => cont (g, p', add_env id dv e', k')
| _ => raise_p pc "IMPOSSIBLE: Ret op in non-return configuration"
end
```

interpreter returns
an interaction tree
with "LLVM" effects.
LLVMTrace := itree IO

```
| Inst insn => (* instruction *)
do pc_next ← trywith "no fallthrough instruction" (incr_pc CFG pc);
match (pt pc), insn with
```

Extract to executable
interpreter (Ocaml).

```
| IIId id, INSTR_Op op =>
'dv ← eval_op g e op;
cont (g, pc_next, add_env id dv e, k)
```

The interpreter
"calls out" to the memory
model by generating
visible effects...

```
| IIId id, INSTR_Alloca i _ _ =>
Trace.Vis (Alloca (eval_typ t))
(λ (a:dvalue) > cont (g, pc_next, add_env id a e, k))
```

```
| IIId id, INSTR_Load _ t (u,ptr) _ =>
'dv ← eval_exp (Some (eval_typ u)) ptr;
Trace.Vis (Load (eval_typ t) dv)
(λ dv > cont (g, pc_next, add_env id dv e, k))
```

Interactive Theorem Proving

```
Theorem block_fusion_cfg_correct :  
  ∀ (G : cfg dtyp),  
    wf_cfg G →  
    ⟦ G ⟧cfg ≈ ⟦ block_fusion_cfg G ⟧cfg.  
  
Proof.  
  intros G [WF1 WF2].  
  unfold denote_cfg.  
  simpl bind.  
  unfold block_fusion_cfg.  
  destruct (block_fusion G.(blk)) as [H1 H2].  
  - break_match_goal; reflexivity.  
  simpl.  
  apply Bool.orb_false_elim in H1.  
  unfold Eqv.eqv_dec in *.
```

In Coq, one can state Lemmas just as easily as any other kind of function.

You can prove those lemmas interactively. Coq checks each step as you do it.

Comparing Behaviors

- Consider two programs P_1 and P_2 possibly in different languages.
 - e.g., P_1 is an Oat program, P_2 is its compilation to LL
- The semantics of the languages associate to each program a set of observable behaviors:

$$\mathcal{B}(P) \text{ and } \mathcal{B}(P')$$

- Note: $|\mathcal{B}(P)| = 1$ if P is deterministic, > 1 otherwise

What is Observable?

- For C-like languages:

observable behavior ::=

- | terminates(st) (i.e. observe the final state)
- | diverges
- | goeswrong

- For pure functional languages:

observable behavior ::=

- | terminates(v) (i.e. observe the final value)
- | diverges
- | goeswrong

What about I/O?

- Add a *trace* of input-output events performed:

coind.	$t ::= [] \mid e :: t$ $T ::= [] \mid e :: T$	(finite traces) (finite and infinite traces)
--------	--	---

observable behavior ::=

terminates(t , st)	(end in state st after trace t)
diverges(T)	(loop, producing trace T)
goeswrong(t)	

Examples

- P1:

`print(1); / st` \Rightarrow `terminates(out(1)::[],st)`

- P2:

`print(1); print(2); / st` \Rightarrow `terminates(out(1)::out(2)::[],st)`

- P3:

`WHILE true DO print(1) END / st` \Rightarrow `diverges(out(1)::out(1)::...)`

- So $\mathfrak{B}(P1) \neq \mathfrak{B}(P2) \neq \mathfrak{B}(P3)$

Bisimulation

- Two programs P1 and P2 are bisimilar whenever:

$$\mathfrak{B}(P1) = \mathfrak{B}(P2)$$

- The two programs are completely indistinguishable.
- But... this is often too strong in practice.

Compilation Reduces Nondeterminism

- Some languages (like C) have underspecified behaviors:
 - Example: order of evaluation of expressions $f() + g()$

- Concurrent programs often permit nondeterminism
 - Classic optimizations can reduce this nondeterminism
 - Example:

$a := x + 1; b := x + 1 \quad || \quad x := x+1$

vs.

$a := x + 1; b := a \quad || \quad x := x+1$

- LLVM explicitly allows nondeterminism:
 - undef values (not part of LLVM lite)
 - see the discussion later

Backward Simulation

- Program P2 can exhibit fewer behaviors than P1:

$$\mathcal{B}(P1) \supseteq \mathcal{B}(P2)$$

- All of the behaviors of P2 are permitted by P1, though some of them may have been eliminated.
- Also called *refinement*.

What about goes wrong?

- Compilers often translate away bad behaviors.

$x := 1/y ; x := 42$ (divide by 0 error)	vs.	$x := 42$ (always terminates)
---	-----	----------------------------------

- Justifications:
 - Compiled program does not “go wrong” because the program type checks or is otherwise formally verified
 - Or just “garbage in/garbage out”

Safe Backwards Simulation

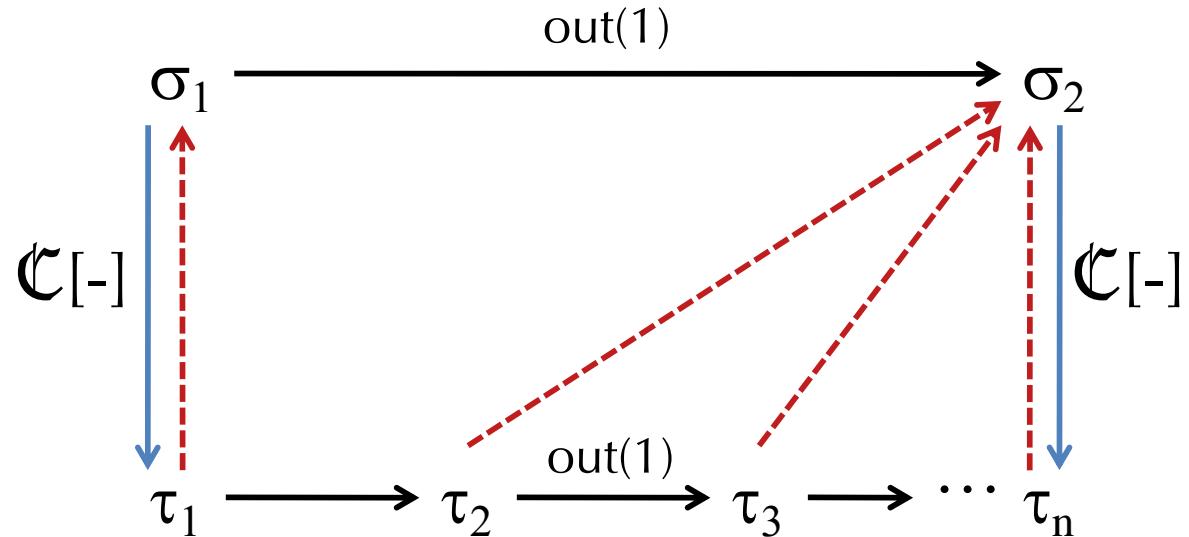
- Only require the compiled program's behaviors to agree if the source program could not go wrong:

$$\text{goeswrong}(t) \notin \mathcal{B}(P_1) \quad \Rightarrow \quad \mathcal{B}(P_1) \supseteq \mathcal{B}(P_2)$$

- Idea: let S be the *functional specification* of the program:
A set of behaviors not containing $\text{goeswrong}(t)$.
 - A program P satisfies the spec if $\mathcal{B}(P) \subseteq S$
- Lemma: If P_2 is a safe backwards simulation of P_1 and P_1 satisfies the spec, then P_2 does too.

Building Backward Simulations

Source:



Idea: The event trace along a (target) sequence of steps originating from a compiled program must correspond to some source sequence.

Tricky parts:

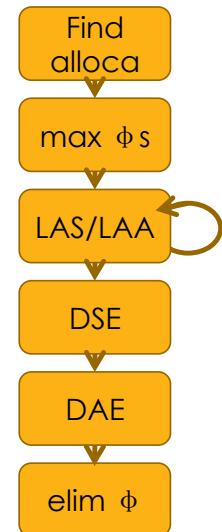
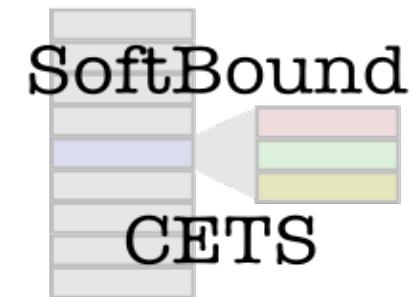
- Must consider all possible target steps
- If the compiler uses many target steps for once source step, we have to invent some way of relating the intermediate states to the source.
- the compilation function goes the wrong way to help!

So What?

- Find bugs in the existing LLVM infrastructure
 - thinking hard about corner cases while formalizing is a good way to find real bugs
 - identify inconsistent assumptions on the LLVM compiler
- Automated Tests against other implementations
 - e.g., integrate with Csmith
- Formally validate program transformations
 - is a particular optimization correct?
 - improve confidence in novel program transformations
- Eventually... verify compiler front ends and/or back ends
 - to obtain a fully-verified CompCert-like compiler

VELLVM [Previous Results]

- Verified **SoftBound**
 - Memory Safety
- Verified **mem2reg**
 - Register promotion, defined in terms of a stack of "micro-optimizations"
- Verified **dominator analysis**
 - Cooper-Harvey-Kennedy Algorithm
- Better memory models
 - ptrtoint casts
 - modular formalization



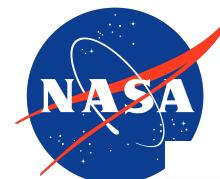
Can it Scale?

- Use of theorem proving to verify “real” software is still considered to be the bleeding edge of research.
- **CompCert** – fully verified C compiler
Leroy, INRIA
- **Vellvm** – formalized LLVM IR
Zdancewic, Penn
- **Ynot** – verified DBMS, web services
Morrisett, Harvard
- **Verified Software Toolchain**
Appel, Princeton
- **Bedrock** – web programming, packet filters
Chlipala, MIT
- **CertiKOS** – certified OS kernel
Shao & Ford, Yale
- **CakeML** – certified compiler
- **SEL4** – certified secure OS microkernel
- **Kami** – verified RISCV architecture
- **DaisyNSF** – verified NFS file system
- ...



AIRBUS

Microsoft



aws

Google Meta

 **BEDROCK**
Systems Inc

| galois |

Formal Methods for Blockchain

Academic Work:

A Survey of Smart Contract Formal Specification and Verification
[Tolmach, et al. 2021]



C E R T I K



Uses deep spec
results



CARDANO

tz Tezos

Where next?

- Proof engineering is still nascent
 - automation, scale, maintenance
 - software engineering++
 - new theory needed: dealing with equality
- Verification is still hard
 - labor intensive, difficult, \$\$\$\$
- Deep Specifications
 - what are the principles?
 - compositionality?
- Real-time, cyberphysical,...

What have we learned?

Where else is it applicable?

What next?

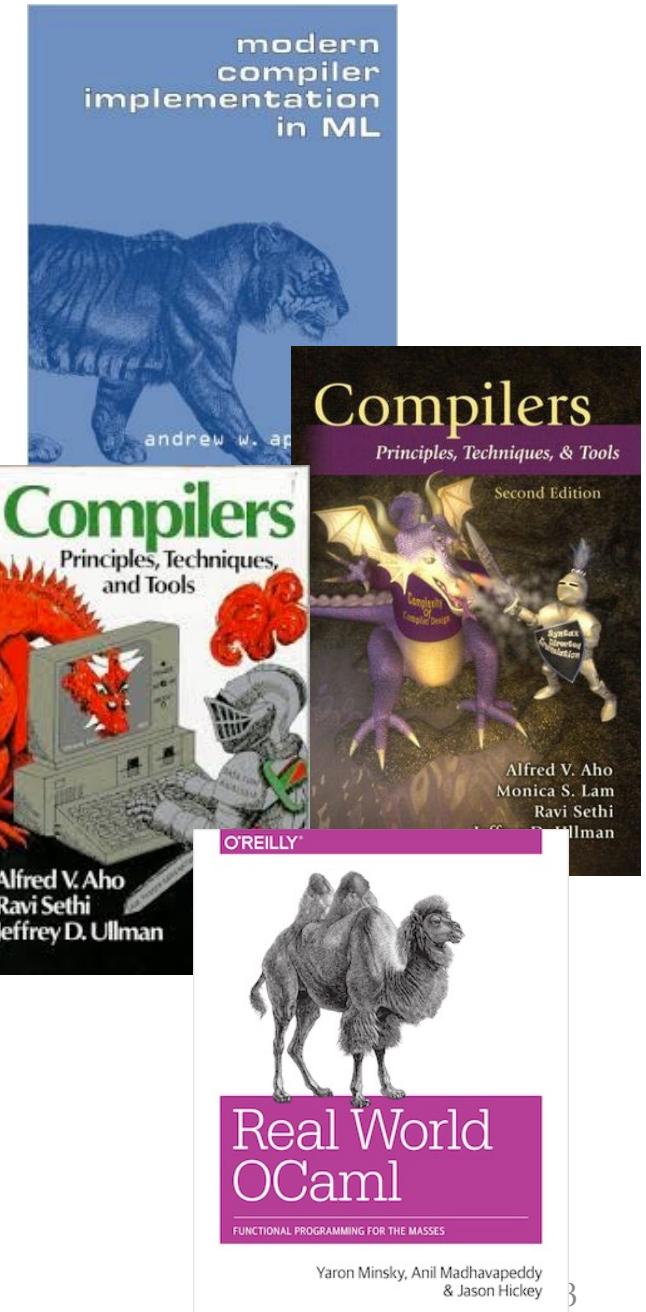
COURSE WRAP-UP

Final Exam

- Will mostly cover material since the midterm
 - Starting from Lecture 14
 - Lambda calculus / closure conversion
 - Scope / Typechecking / Inference Rules
 - Objects, inheritance, types, implementation of dynamic dispatch (de-emphasized, since we didn't cover it thoroughly)
 - Basic optimizations
 - Dataflow analysis (forward vs. backward, fixpoint computations, etc.)
 - Liveness
 - Graph-coloring Register Allocation
 - Control flow analysis
 - Loops, dominator trees
- One, letter-sized, double-sided, hand-written “cheat sheet”

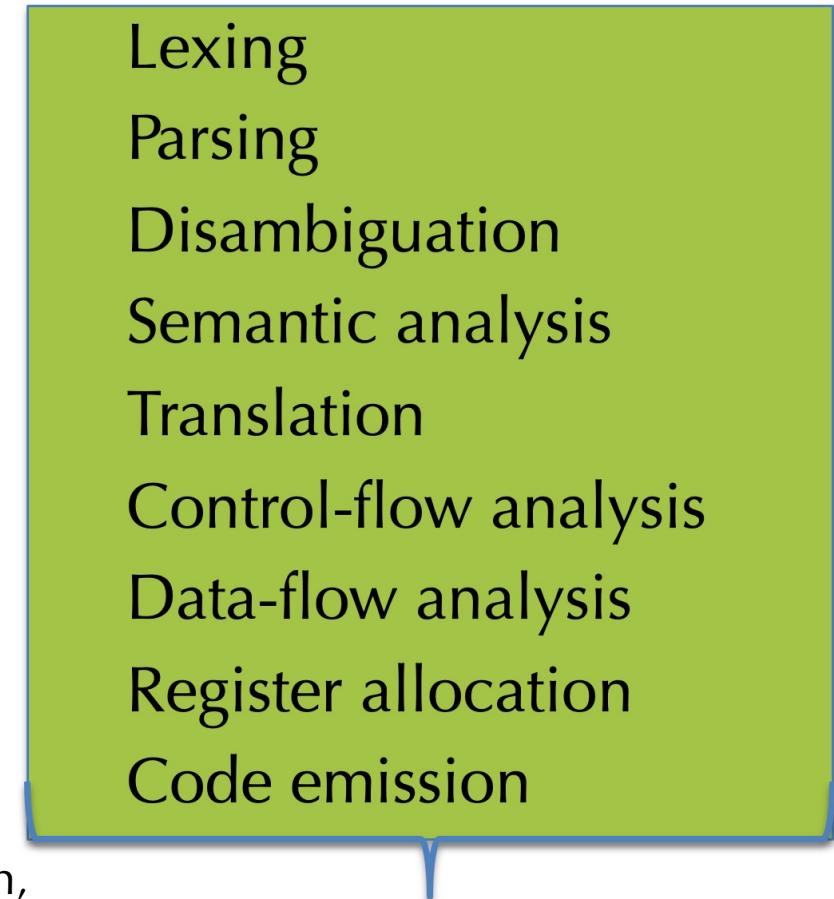
Why CIS 341?

- You will learn:
 - Practical applications of theory
 - Parsing
 - How high-level languages are implemented in machine language
 - (A subset of) Intel x86 architecture
 - A deeper understanding of code
 - A little about programming language semantics
 - Functional programming in OCaml
 - How to manipulate complex data structures
 - How to be a better programmer
- Did we meet these goals?



Stuff we didn't Cover

- We skipped stuff at every level...
- Concrete syntax/parsing:
 - Much more to the theory of parsing...
 $LR(*)$
 - Good syntax is art, not science!
- Source language features:
 - Exceptions, advanced type systems, type inference, concurrency
- Intermediate languages:
 - Intermediate language design, bytecode, bytecode interpreters, just-in-time compilation (JIT)
- Compilation:
 - Continuation-passing transformation, efficient representations, scalability
- Optimization:
 - Scientific computing, cache optimization, instruction selection/optimization
- Runtime support:
 - memory management, garbage collection



Lexing
Parsing
Disambiguation
Semantic analysis
Translation
Control-flow analysis
Data-flow analysis
Register allocation
Code emission

Compiler Passes

Related Courses

- CIS 500: Software Foundations
 - Prof. Pierce
 - Theoretical course about functional programming, proving program properties, type systems, lambda calculus. Uses the theorem prover Coq.
- CIS 501: Computer Architecture
 - Prof. Devietti
 - 371++: pipelining, caches, VM, superscalar, multicore,...
- CIS 547: Software Analysis
 - Prof. Naik
 - LLVM IR + program analysis
- CIS 552: Advanced Programming
 - Prof. Weirich
 - Advanced functional programming in Haskell, including generic programming, metaprogramming, embedded languages, cool tricks with fancy type systems
- CIS 670: Special topics in programming languages

Where to go from here?

- Conferences (proceedings available on the web):
 - Programming Language Design and Implementation (PLDI)
 - Principles of Programming Languages (POPL)
 - Object Oriented Programming Systems, Languages & Applications (OOPSLA)
 - International Conference on Functional Programming (ICFP)
 - European Symposium on Programming (ESOP)
 - ...
- Technologies / Open Source Projects
 - Yacc, lex, bison, flex, ...
 - LLVM – low level virtual machine
 - Java virtual machine (JVM), Microsoft's Common Language Runtime (CLR)
 - Languages: OCaml, F#, Haskell, Scala, Go, Rust, ...?

Where else is this stuff applicable?

- General programming
 - In C/C++, better understanding of how the compiler works can help you generate better code.
 - Ability to read assembly output from compiler
 - Experience with functional programming can give you different ways to think about how to solve a problem
- Writing domain specific languages
 - lex/yacc very useful for little utilities
 - understanding abstract syntax and interpretation
- Understanding hardware/software interface
 - Different devices have different instruction sets, programming models

Thanks!

- To the TAs: Stephen, Lef, and Sumanth
- To *you* for taking the class!
- How can I improve the course?
 - Let me know in course evaluations!