

CS 516: COMPILERS

Lecture 3

Topics

- Introduction to X86lite.

Materials

- `lec03.zip` (Homework infrastructure)

CS 516 Announcements

- HW1: Hellocaml!
 - is due *Thursday night* at 11:59:59pm.
- HW2: X86lite Simulator
 - Available on the course web pages.
 - Due: *next Thursday* Feb. 9th at 11:59pm
 - Pair-programming project
 - Simulator for x86 Assembly subset
- **Office Hours**
 - Ben - Wednesdays 3-5pm and Thursdays 4-5pm
 - Vidya - Mondays 1pm-3pm
 - me - Mondays 3-4:30m

HW Infrastructure

- HW2: X86lite Simulator (HW2 and beyond)
 - Use an ocaml infrastructure

```
unzip lec03.zip  
cd lec03/  
dune build  
bin/main.exe
```

(Finish lec02 first.)

Demo Interpreter/Compiler

1. Interpreter – `simple.ml` (`simple-soln.ml`)
2. Compiler – `translate.ml`

```
cd simple
dune build
dune exec bin/simple.exe




dune utop
utop # #use "bin/simple.ml";;
utop # let s' = (interpret_cmd init_state factorial);;
utop # lookup s' "ANS";;
```

The target architecture for CS 516

X86LITE

Intel's X86 Architecture

- 1978: Intel introduces 8086
- 1982: 80186, 80286
- 1985: 80386
- 1989: 80486 (100MHz, 1 μ m)
- 1993: Pentium
- 1995: Pentium Pro
- 1997: Pentium II/III
- 2000: Pentium 4
- 2003: Pentium M, Intel Core
- 2006: Intel Core 2
- 2008: Intel Core i3/i5/i7
- 2011: SandyBridge / IvyBridge
- 2013: Haswell
- 2014: Broadwell
- 2015: Skylake (4.2GHz, 14nm)
- 2016: Xeon Phi
- AMD has a parallel line of processors




X86 Evolution & Moore's Law

Moore's Law: The number of transistors on microchips doubles every two years

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important for other aspects of technological progress in computing – such as processing speed or the price of computers.

Our World
in Data

Transistor count

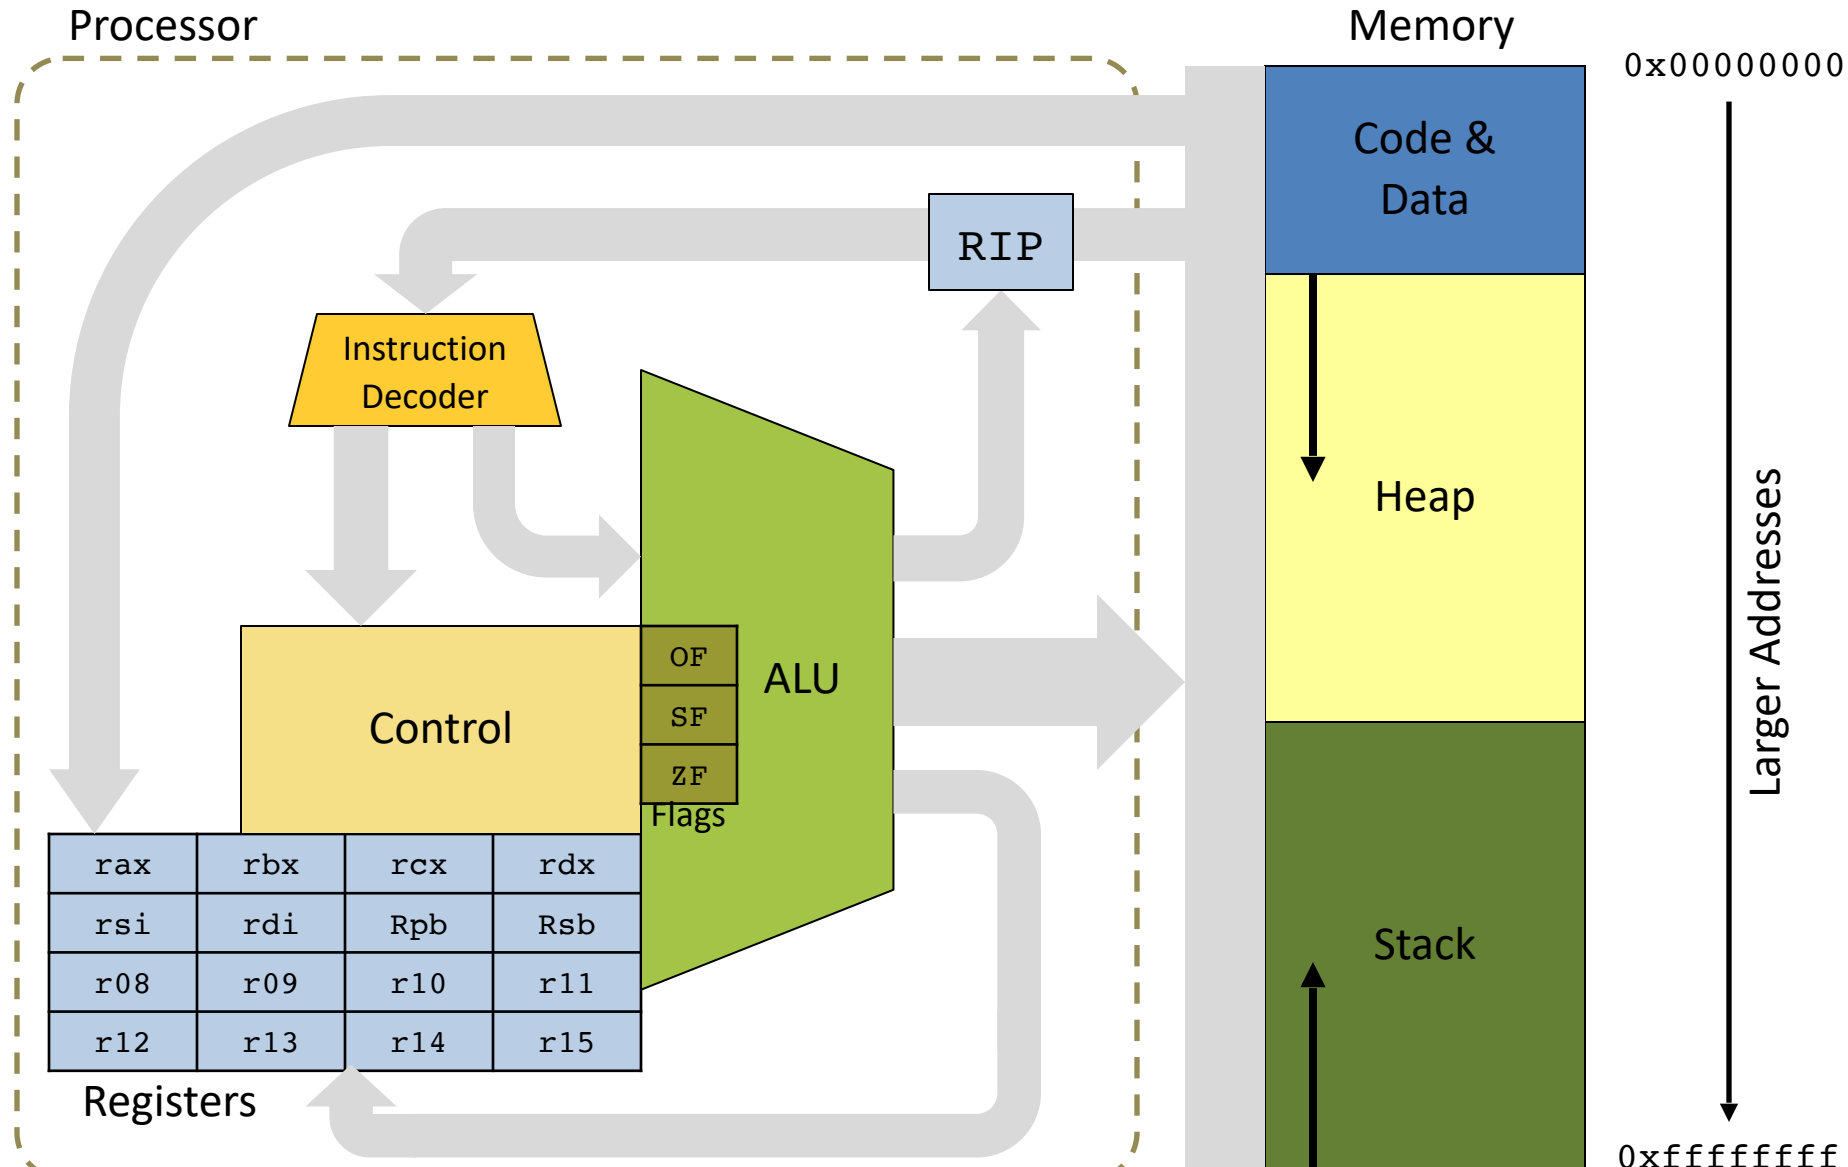


Data source: Wikipedia (wikipedia.org/wiki/Transistor_count)

X86 vs. X86lite

- X86 assembly is *very* complicated:
 - 8-, 16-, 32-, 64-bit values + floating points, etc.
 - Intel 64 and IA 32 architectures have a *huge* number of functions
 - “CISC” complex instructions
 - Machine code: instructions range in size from 1 byte to 17 bytes
 - Lots of hold-over design decisions for backwards compatibility
 - Hard to understand, there is a large book about optimizations at just the instruction-selection level
- X86lite is a *very* simple subset of X86:
 - Only 64 bit signed integers (no floating point, no 16bit, no ...)
 - Only about 20 instructions
 - Sufficient as a target language for general-purpose computing

X86 Schematic



X86lite Machine State: Registers

- Register File: 16 64-bit registers
 - **rax** general purpose accumulator
 - **rbx** base register, pointer to data
 - **rcx** counter register for strings & loops
 - **rdx** data register for I/O
 - **rsi** pointer register, string source register
 - **rdi** pointer register, string destination register
 - **rbp** base pointer, points to the stack frame
 - **rsp** stack pointer, points to the top of the stack
 - **R08–r15** general purpose registers
- **rip** a “virtual” register, points to the current instruction
 - **rip** is manipulated only indirectly via jumps and return.

Simplest instruction: mov

- `movq SRC, DEST` copy SRC into DEST
- Here, DEST and SRC are *operands*
- DEST is treated as a *location*
 - A location can be a register or a memory address
- SRC is treated as a *value*
 - A value is the *contents* of a register or memory address
 - A value can also be an *immediate* (constant) or a label
- `movq $4, %rax` // move the 64-bit immediate value 4 into rax
- `movq %rbx, %rax` // move the contents of rbx into rax

A Note About Instruction Syntax

- X86 presented in *two* common syntax formats
- AT&T notation: source *before* destination
 - Prevalent in the Unix/Mac ecosystems
 - Immediate values prefixed with '\$'
 - Registers prefixed with '%'
 - Mnemonic suffixes: `movq` vs. `mov`
 - `q` = quadword (4 words)
 - `l` = long (2 words)
 - `w` = word
 - `b` = byte
- Intel notation: destination *before* source
 - Used in the Intel specification / manuals
 - Prevalent in the Windows ecosystem
 - Instruction variant determined by register name

```
movq $5, %rax
movl $5, %eax
```

src dest

```
mov rax, 5
mov eax, 5
```

dest src

A Note About Instruction Syntax

- X86 presented in *two* common syntax formats

- AT&T notation: source *before* destination

- Prevalent in the Unix/Mac ecosystems
- Immediate values prefixed with '\$'
- Registers prefixed with '%'
- Mnemonic suffixes: `movq` vs. `mov`
 - `q` = quadword (4 words)
 - `l` = long (2 words)
 - `w` = word
 - `b` = byte

```
movq $5, %rax
```

```
movl $5, %eax
```

src dest

Note: X86lite uses the AT&T notation and the 64-bit only version of the instructions and registers.

- Intel notation: destination *before* source
 - Used in the Intel specification / manuals
 - Prevalent in the Windows ecosystem
 - Instruction variant determined by register name

```
mov rax, 5
```

```
mov eax, 5
```

dest src

X86lite Arithmetic instructions

- `negq DEST` two's complement negation
- `addq SRC, DEST` $DEST \leftarrow DEST + SRC$
- `subq SRC, DEST` $DEST \leftarrow DEST - SRC$
- `imulq SRC, Reg` $Reg \leftarrow Reg * SRC$ (truncated 128-bit mult.)

Examples as written in:

```
addq %rbx, %rax    // rax ← rax + rbx
subq $4, rsp       // rsp ← rsp - 4
```

- Note: `Reg` (in `imulq`) must be a register, not a memory address

X86lite Logic/Bit manipulation Operations

- `notq DEST` logical negation
- `andq SRC, DEST` $DEST \leftarrow DEST \ \&\& \ SRC$
- `orq SRC, DEST` $DEST \leftarrow DEST \ || \ SRC$
- `xorq SRC, DEST` $DEST \leftarrow DEST \ xor \ SRC$

- `sarq Amt, DEST` $DEST \leftarrow DEST \ >> \ amt$ (arithmetic shift right)
- `shlq Amt, DEST` $DEST \leftarrow DEST \ << \ amt$ (arithmetic shift left)
- `shrq Amt, DEST` $DEST \leftarrow DEST \ >>> \ amt$ (bitwise shift right)


X86lite Logic/Bit manipulation Operations

- `notq DEST` logical negation
- `andq SRC, DEST` $DEST \leftarrow DEST \&\& SRC$
- `orq SRC, DEST` $DEST \leftarrow DEST || SRC$
- `xorq SRC, DEST` $DEST \leftarrow DEST \text{ xor } SRC$

- `sarq Amt, DEST` $DEST \leftarrow DEST \gg amt$ (arithmetic shift right)
- `shlq Amt, DEST` $DEST \leftarrow DEST \ll amt$ (arithmetic shift left)
- `shrq Amt, DEST` $DEST \leftarrow DEST \ggg amt$ (bitwise shift right)

Reminder:

- In arithmetic shift right, the MSB is replicated (diagram).
- In arithmetic shift left, the LSB is filled with a 0.
- In (logical) shift right, the MSB is 0.



X86 Operands

- Operands are the values operated on by the assembly instructions
 - **Imm** 64-bit literal signed integer “immediate”
 - **Lbl** a “label” representing a machine address
the assembler/linker/loader resolve labels
 - **Reg** One of the 16 registers, the value of a register is
its contents
 - **Ind** [base:Reg][index:Reg,scale:int32][disp]
machine address (see next slide)

X86 Addressing


- In general, there are three components of an indirect address
 - **Base:** a machine address stored in a register
 - **Index * scale:** a variable offset from the base
 - **Disp:** a constant offset (displacement) from the base
- $\text{addr(ind)} = \text{Base} + [\text{Index} * \text{scale}] + \text{Disp}$
 - When used as a *location*, ind denotes the address addr(ind)
 - When used as a *value*, ind denotes $\text{Mem}[\text{addr(ind)}]$, the contents of the memory address
- Example: $-4(\%rsp)$ denotes address: $rsp - 4$
- Example: $(\%rax, \%rcx, 4)$ denotes address: $rax + 4 * rcx$
- Example: $12(\%rax, \%rcx, 4)$ denotes address: $rax + 4 * rcx + 12$
- Note: Index cannot be `rsp`

X86 Addressing

- In general, there are three components of an indirect address
 - **Base:** a machine address stored in a register
 - **Index * scale:** a variable offset from the base
 - **Disp:** a constant offset (displacement) from the base
- $\text{addr(ind)} = \text{Base} + [\text{Index} * \text{scale}] + \text{Disp}$
 - When used as a *location*, ind denotes the address addr(ind)
 - When used as a *value*, ind denotes $\text{Mem}[\text{addr(ind)}]$, the contents of the memory address
- Example: $-4(\%rsp)$ denotes address: $rsp - 4$
- Example: $(\%rax, \%rcx, 4)$ denotes address: $rax + 4 * rcx$
- Example: $12(\%rax, \%rcx, 4)$ denotes address: $rax + 4 * rcx + 12$
- Note: Index cannot be `rsp`

Note: X86lite does not need this full generality. It does not use `index * scale`.

X86lite Memory Model



- 2^{64} bytes numbered `0x00000000` through `0xffffffff`.
- Treated as 64-bit (8-byte) quadwords.
- Therefore: legal X86lite memory addresses consist of *64-bit, quadword-aligned pointers*.
 - All memory addresses are evenly divisible by 8

X86lite Memory Model

- Useful instruction:

`leaq Ind, DEST` $DEST \leftarrow \text{addr}(\text{Ind})$ loads a pointer into DEST

- By convention, there is a stack that grows from high addresses to low addresses
- The register `rsp` points to the top of the stack
 - `pushq SRC` $rsp \leftarrow rps - 8; \text{Mem}[rsp] \leftarrow SRC$
 - `popq DEST` $DEST \leftarrow \text{Mem}[rsp]; rsp \leftarrow rsp + 8$

X86lite State: Condition Flags & Codes

X86lite State: Condition Flags & Codes

- X86 instructions set flags as a side effect

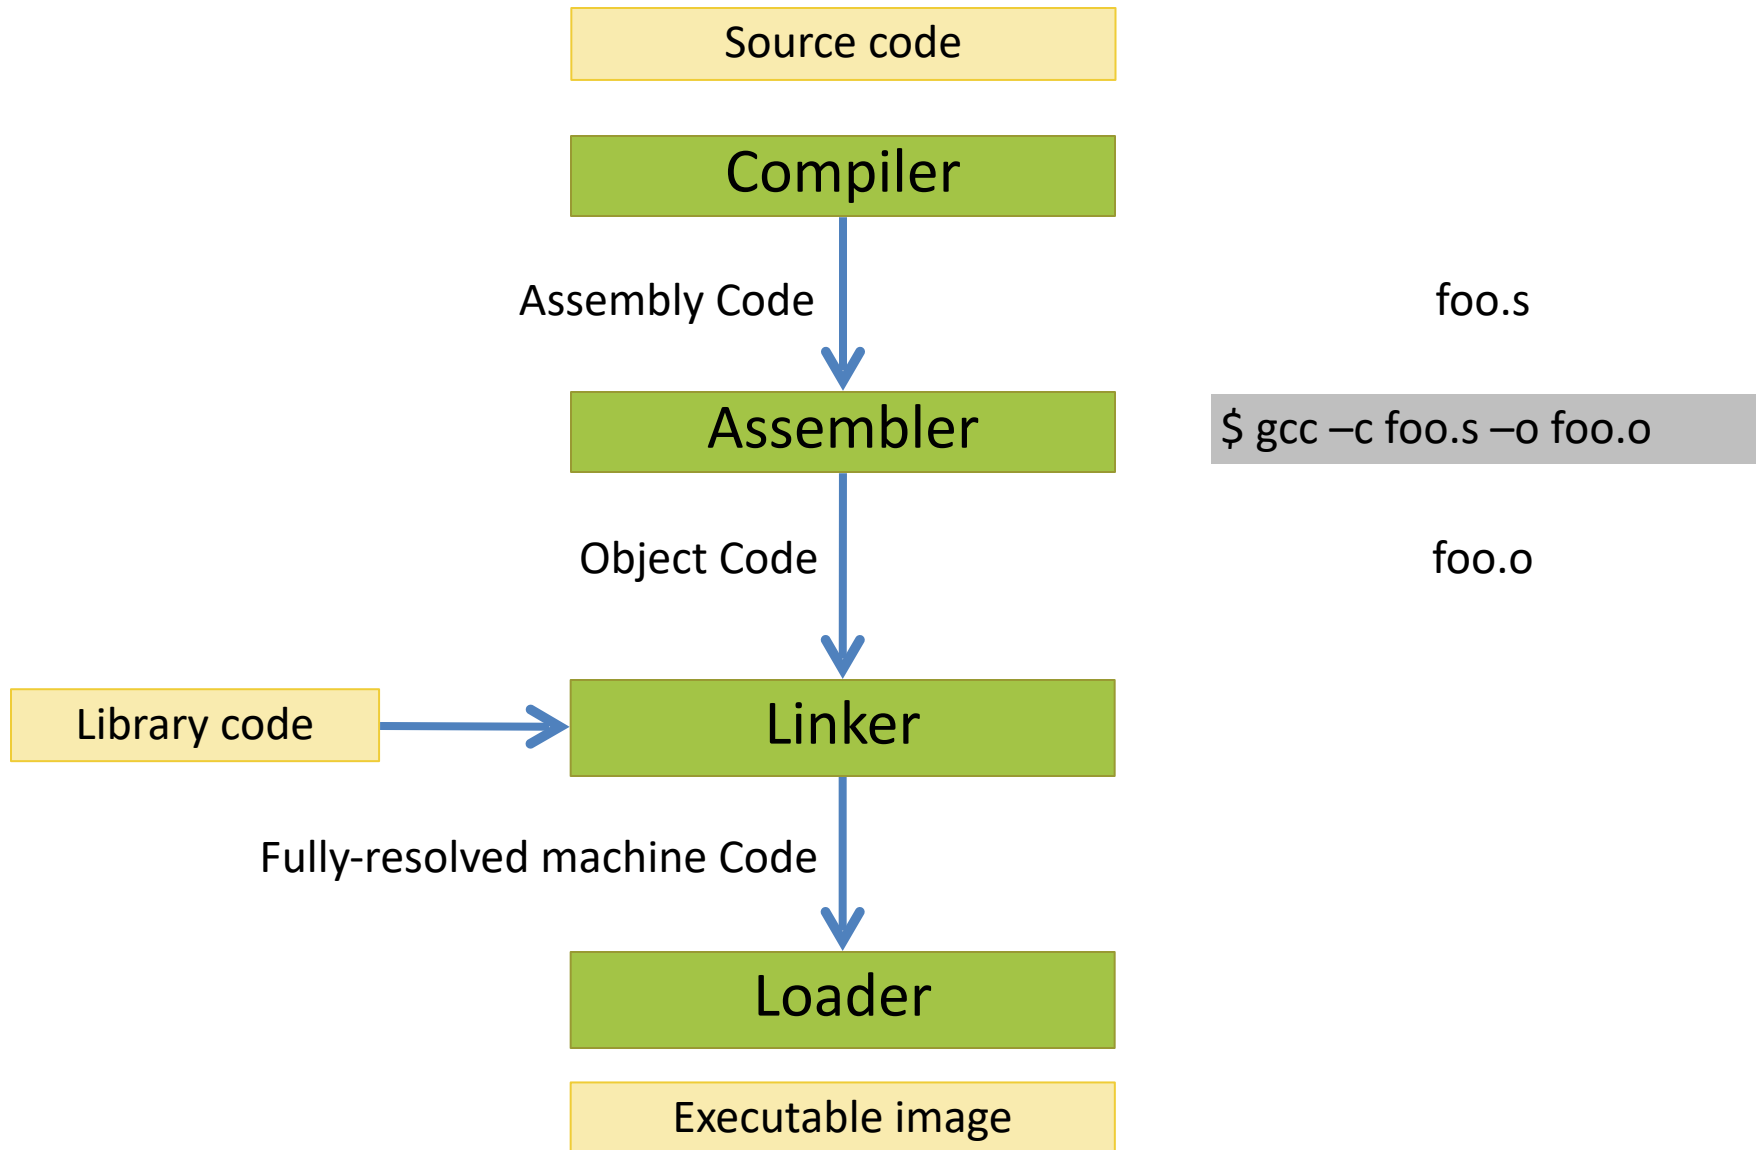
X86lite State: Condition Flags & Codes

- X86 instructions set flags as a side effect
- X86lite has only 3 flags:
 - OF: “**overflow**” set when the result is too big/small to fit in 64-bit reg.
 - SF: “**sign**” set to the sign of the result (0=positive, 1 = negative)
 - ZF: “**zero**” set when the result is 0

X86lite State: Condition Flags & Codes

- X86 instructions set flags as a side effect
- X86lite has only 3 flags:
 - **OF**: “**overflow**” set when the result is too big/small to fit in 64-bit reg.
 - **SF**: “**sign**” set to the sign of the result (0=positive, 1 = negative)
 - **ZF**: “**zero**” set when the result is 0
- From these flags, we can define *Condition Codes*
 - To compare SRC1 and SRC2, compute SRC1 – SRC2 to set the flags
 - **e** equality holds when **ZF** is set
 - **ne** inequality holds when (**not ZF**)
 - **g** greater than holds when **not (SF <> OF || ZF)**
 - **l** less than holds when **SF <> OF**
 - Equivalently: $((\text{SF} \ \&\& \ \text{not } \text{OF}) \ || \ (\text{not } \text{SF} \ \&\& \ \text{OF}))$
 - **ge** greater or equal holds when **not (SF <> OF)**
 - **le** than or equal holds when **SF <> OF or ZF**

Compilation & Execution



Code Blocks & Labels


- X86 assembly code is organized into *labeled blocks*:

```
label1:  
    <instruction>  
    <instruction>  
    ...  
    <instruction>  
  
label2:  
    <instruction>  
    <instruction>  
    ...  
    <instruction>
```

- Labels indicate code locations that can be jump targets (either through conditional branch instructions or function calls).
- Labels are translated away by the linker and loader – instructions live in the heap in the “code segment”
- An X86 program begins executing at a designated code label (usually “main”).

Conditional Instructions

- First compare, then jump based on comparison
- `cmpq SRC1, SRC2` Compute $SRC2 - SRC1$, set condition flags
 - **ZF=1** if $SRC1 = SRC2$
 - **SF=1** if $(SRC2 - SRC1) < 0$ (signed)
 - **OF=1** if two's complement (signed overflow)
- Example:
`cmpq %rcx, %rax` Compare `rax` to `ecx`
`je __true1b1` If `rax = rcx` then jump to `__true1b1`



Conditional Instructions


- `jCC SRC` `rip` \leftarrow if CC then SRC else fallthrough

- Example:

```
cmpq %rcx, %rax
je __true1b1
```

Compare rax to ecx

If `rax = rcx` then jump to `__true1b1`



Conditional Instructions

- `jCC SRC`

`rip` \leftarrow if CC then SRC else fallthrough

Instruction	Condition	Description
<code>jmp target</code>	1	Unconditional
<code>je target</code>	ZF	Equal / Zero
<code>jne target</code>	\sim ZF	Not Equal / Not Zero
<code>js target</code>	SF	Negative
<code>jns target</code>	\sim SF	Nonnegative
<code>jg target</code>	$\sim (SF \wedge OF) \ \& \ \sim ZF$	Greater (Signed)
<code>jge target</code>	$\sim (SF \wedge OF)$	Greater or Equal (Signed)
<code>jl target</code>	$(SF \wedge OF)$	Less (Signed)
<code>jle target</code>	$(SF \wedge OF) \mid ZF$	Less or Equal (Signed)
<code>ja target</code>	$\sim CF \ \& \ \sim ZF$	Above (unsigned ">")
<code>jb target</code>	CF	Below (unsigned "<")

- Example:

```
cmpq %rcx, %rax
```

Compare rax to ecx

```
je __true1b1
```

If rax = rcx then jump to __true1b1

CF	Carry Flag	ZF	Zero Flag	SF	Sign Flag	OF	Overflow Flag
----	------------	----	-----------	----	-----------	----	---------------

Conditional Instructions

Reading Condition Codes

Register	Use(s)
%rdi	1 st argument (x)
%rsi	2 nd argument (y)
%rax	return value

❖ set* Instructions

- Set a low-order byte to 0 or 1 based on condition codes
- Operand is byte register (e.g. al, dl) or a byte in memory
- Do not alter remaining bytes in register
 - Typically use `movzbl` (zero-extended mov) to finish job

```
int gt(long x, long y)
{
    return x > y;
}
```

```
cmpq    %rsi, %rdi    # Compare x:y
setg     %al           # Set when >
movzbl   %al, %eax     # Zero rest of %rax
ret
```


Conditional Instructions

- Reading the condition code
- `setbCC DEST` `DEST's lower byte` \leftarrow if `CC` then 1 else 0

Instruction	Condition	Description
sete <i>dst</i>	ZF	Equal / Zero
setne <i>dst</i>	\sim ZF	Not Equal / Not Zero
sets <i>dst</i>	SF	Negative
setns <i>dst</i>	\sim SF	Nonnegative
setg <i>dst</i>	$\sim (SF \wedge OF) \ \& \ \sim ZF$	Greater (Signed)
setge <i>dst</i>	$\sim (SF \wedge OF)$	Greater or Equal (Signed)
setl <i>dst</i>	$(SF \wedge OF)$	Less (Signed)
setle <i>dst</i>	$(SF \wedge OF) \mid ZF$	Less or Equal (Signed)
seta <i>dst</i>	$\sim CF \ \& \ \sim ZF$	Above (unsigned ">")
setb <i>dst</i>	CF	Below (unsigned "<")

CF Carry Flag **ZF** Zero Flag **SF** Sign Flag **OF** Overflow Flag

Jumps, Call and Return

- `jmp SRC` `rip` \leftarrow SRC Jump to location in SRC
- `callq SRC` Push `rip`; `rip` \leftarrow SRC
 - Call a procedure: Push the program counter to the stack (decrementing `rsp`) and then jump to the machine instruction at the address given by SRC.
- `retq` Pop into `rip`
 - Return from a procedure: Pop the current top of the stack into `rip` (incrementing `rsp`).
 - This instruction effectively jumps to the address at the top of the stack