

Lecture 20

# CIS 341: COMPILERS

# Announcements

- HW5: Oat v. 2.0
  - records, function pointers, type checking, array-bounds checks, etc.
  - typechecker & safety
  - Due: Wednesday, April 13<sup>th</sup>



# COMPILING CLASSES AND OBJECTS

# Code Generation for Objects

- Classes:
  - Generate data structure types
    - For objects that are instances of the class and for the class tables
  - Generate the class tables for dynamic dispatch
- Methods:
  - Method body code is similar to functions/closures
  - Method calls require *dispatch*
- Fields:
  - Issues are the same as for records
  - Generating access code
- Constructors:
  - Object initialization
- Dynamic Types:
  - Checked downcasts
  - “instanceof” and similar type dispatch

# Compiling Constructors

- Java and C++ classes can declare constructors that create new objects.
  - Initialization code may have parameters supplied to the constructor
  - e.g. `new Color(r,g,b);`
- Modula-3: object constructors take no parameters
  - e.g. `new Color;`
  - Initialization would typically be done in a separate method.
- Constructors are compiled just like static methods, except:
  - The “this” variable is initialized to a newly allocated block of memory big enough to hold D.V. pointer + fields according to object layout
  - Constructor code initializes the fields
    - What methods (if any) are allowed?
  - The D.V. pointer is initialized
    - When? Before/After running the initialization code?

# Compiling Checked Casts

- How do we compile downcast in general? Consider this generalization of Oat's checked cast:

```
if? (t x = exp) { ... } else { ... }
```

- Reason by cases:
  - t must be either null, ref or ref? (can't be just int or bool)
- If t is null:
  - The static type of exp must be ref? for some ref.
  - If exp == null then take the true branch, otherwise take the false branch
- If t is string or t[]:
  - The static type of exp must be the corresponding string? Or t[]?
  - If exp == null take the false branch, otherwise take the true branch
- If t is C:
  - The static type of exp must be D or D? (where C <: D)
  - If exp == null take the false branch, otherwise:
  - emit code to walk up the class hierarchy starting at D, looking for C
  - If found, then take true branch else take false branch
- If t is C?:
  - The static type of exp must be D? (where C <: D)
  - If exp == null take the true branch, otherwise:
  - Emit code to walk up the class hierarchy starting at D, looking for C
  - If found, then take true branch else take false branch

# “Walking up the Class Hierarchy”

- A non-null object pointer refers to an LLVM struct with a type like:

```
%B = type { %_class_B*, i64, i64, i64 }
```

- The first entry of the struct is a pointer to the vtable for Class B
  - This pointer *is* the dynamic type of the object.
  - It will have the value `@vtbl_B`
- The first entry of the class table for B is a pointer to its superclass:

```
@vtbl_B = global %_class_B { %_class_A* @vtbl_A,  
                             void (%B*)* @print_B,  
                             i64 (%A*, %A*)* @blah_A }
```

- Therefore, to find out whether an unknown type X is a subtype of C:
    - Assume C is not Object (ruled out by “silliness” checks for downcast)
- LOOP:
- If `X == @vtbl_Object` then NO, X is not a subtype of C
  - If `X == @vtbl_C` then YES, X is a subtype of C
  - If `X = @vtbl_D`, so set X to `@vtbl_E` where E is D’s parent and goto LOOP



# MULTIPLE INHERITANCE



# Method Dispatch (Single Inheritance)

- Idea: every method has its own small integer index.
- Index is used to look up the method in the dispatch vector.

```
interface A {  
    void foo();  
}
```

Index

0

```
interface B extends A {  
    void bar(int x);  
    void baz();  
}
```

1

2

Inheritance / Subtyping:

C <: B <: A

```
class C implements B {  
    void foo() {...}  
    void bar(int x) {...}  
    void baz() {...}  
    void quux() {...}  
}
```

0

1

2

3

# Multiple Inheritance

- C++: a class may declare more than one superclass.

- Semantic problem: Ambiguity

```
class A { int m(); }  
class B { int m(); }  
class C extends A,B {...}    // which m?
```

- Same problem can happen with fields.
- In C++, fields and methods can be duplicated when such ambiguity arises (though explicit sharing can be declared too)

- Java: a class may implement more than one interface.
  - No semantic ambiguity: if two interfaces contain the same method declaration, then the class will implement a single method

```
interface A { int m(); }  
interface B { int m(); }  
class C implements A,B {int m() {...}}    // only one m
```

# Dispatch Vector Layout Strategy Breaks

	D.V.Index
<pre>interface Shape {</pre>	
<pre>void setCorner(int w, Point p);</pre>	0
<pre>}</pre>	
<pre>interface Color {</pre>	
<pre>float get(int rgb);</pre>	0
<pre>void set(int rgb, float value);</pre>	1
<pre>}</pre>	
<pre>class Blob implements Shape, Color {</pre>	
<pre>void setCorner(int w, Point p) {...}</pre>	0?
<pre>float get(int rgb) {...}</pre>	0?
<pre>void set(int rgb, float value) {...}</pre>	1?
<pre>}</pre>	

# General Approaches

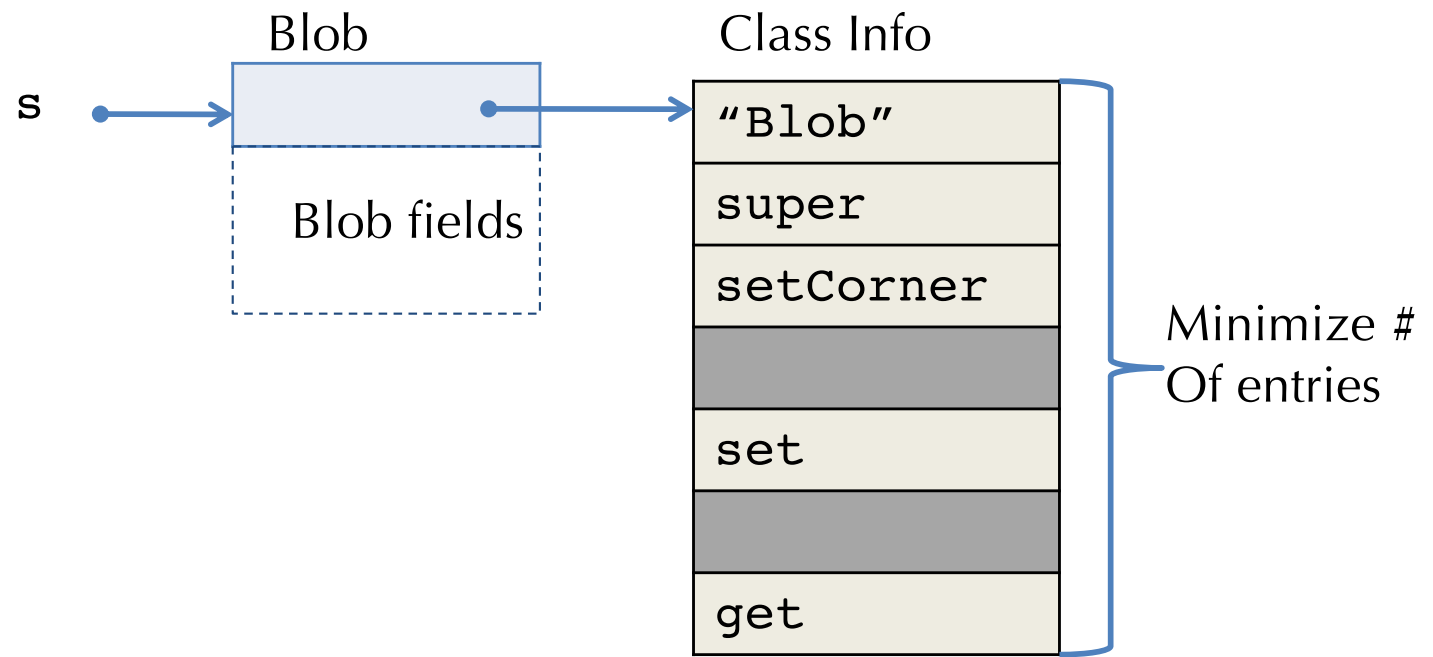
- Can't directly identify methods by position anymore.
- Option 1: Use a level of indirection:
  - Map method identifiers to code pointers (e.g. index by method name)
  - Use a hash table
  - May need to do search up the class hierarchy
- Option 2: Give up separate compilation
  - Use “sparse” dispatch vectors, or binary decision trees
  - Must know then entire class hierarchy
- Option 3: Allow multiple D.V. tables (C++)
  - Choose which D.V. to use based on static type
  - Casting from/to a class may require run-time operations
- Note: many variations on these themes
  - Different Java compilers pick different approaches to options1 and 2...

# Option 2 variant 1: Sparse D.V. Tables

- Give up on separate compilation...
- Now we have access to the whole class hierarchy.
- So: ensure that no two methods in the same class are allocated the same D.V. offset.
  - Allow holes in the D.V. just like the hash table solution
  - Unlike hash table, there is never a conflict!
- Compiler needs to construct the method indices
  - Graph coloring techniques can be used to construct the D.V. layouts in a reasonably efficient way (to minimize size)
  - Finding an optimal solution is NP complete!

# Example Object Layout

- Advantage: Identical dispatch and performance to single-inheritance case
- Disadvantage: Must know entire class hierarchy



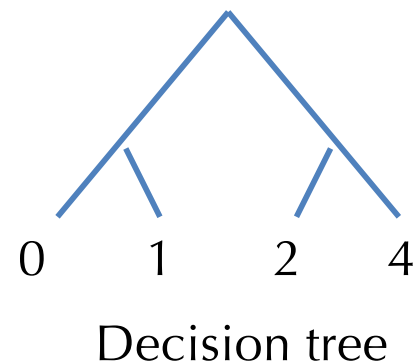
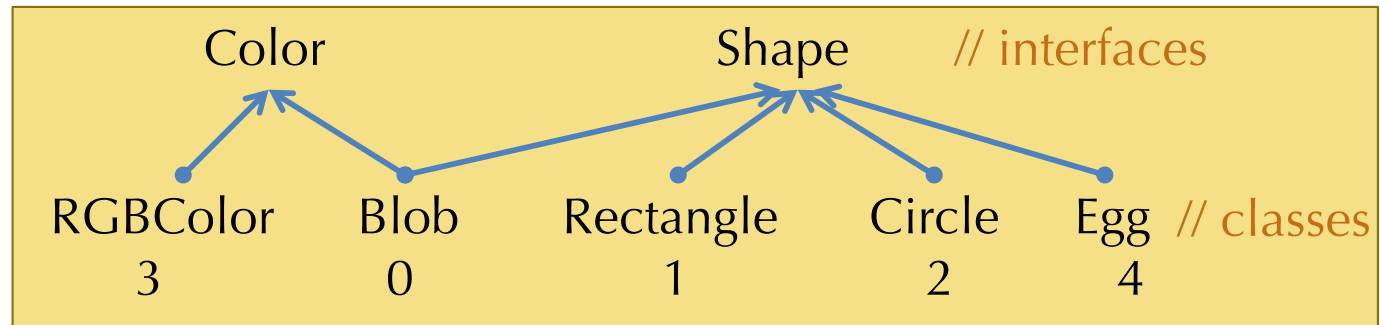
# Option 2 variant 2: Binary Search Trees

- Idea: Use conditional branches not indirect jumps
- Each object has a class index (unique per class) as first word
  - Instead of D.V. pointer (no need for one!)
- Method invocation uses range tests to select among  $n$  possible classes in  $\lg n$  time
  - Direct branches to code at the leaves.

```
Shape x;  
x.SetCorner(...);
```



```
Mov eax, [x]  
Mov ebx, [eax]  
Cmp ebx, 1  
Jle __L1  
Cmp ebx, 2  
Je __CircleSetCorner  
Jmp __EggSetCorner  
__L1:  
Cmp ebx, 0  
Je __BlobSetCorner  
Jmp __RectangleSetCorner
```



# Search Tree Tradeoffs

- Binary decision trees work well if the distribution of classes that may appear at a call site is skewed.
  - Branch prediction hardware eliminates the branch stall of ~10 cycles (on X86)
- Can use profiling to find the common paths for each call site individually
  - Put the common case at the top of the decision tree (so less search)
  - 90%/10% rule of thumb: 90% of the invocations at a call site go to the same class
- Drawbacks:
  - Like sparse D.V.'s you need the whole class hierarchy to know how many leaves you need in the search tree.
  - Indirect jumps can have better performance if there are >2 classes (at most one mispredict)

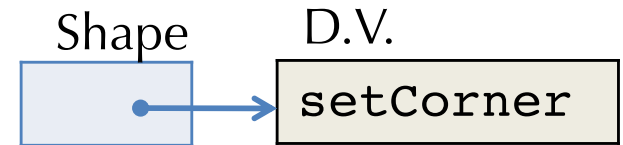


# Option 3: Multiple Dispatch Vectors

- Duplicate the D.V. pointers in the object representation.
- Static type of the object determines which D.V. is used.

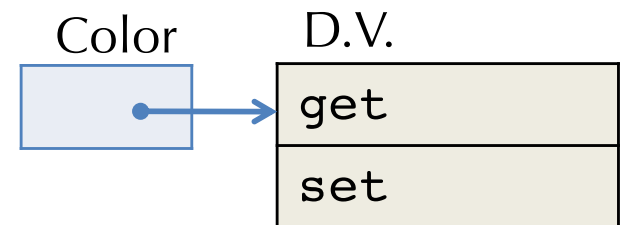
```
interface Shape {  
    void setCorner(int w, Point p);  
}
```

D.V.Index  
0

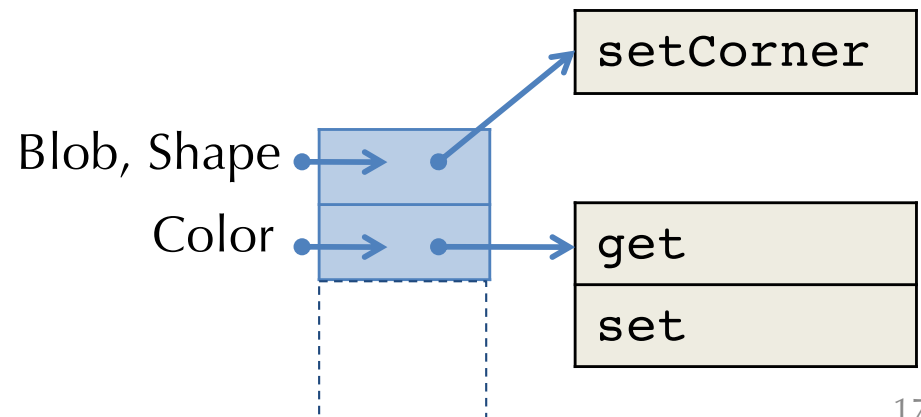


```
interface Color {  
    float get(int rgb);  
    void set(int rgb, float value);  
}
```

0  
1



```
class Blob implements Shape, Color {  
    void setCorner(int w, Point p) {...}  
    float get(int rgb) {...}  
    void set(int rgb, float value) {...}  
}
```



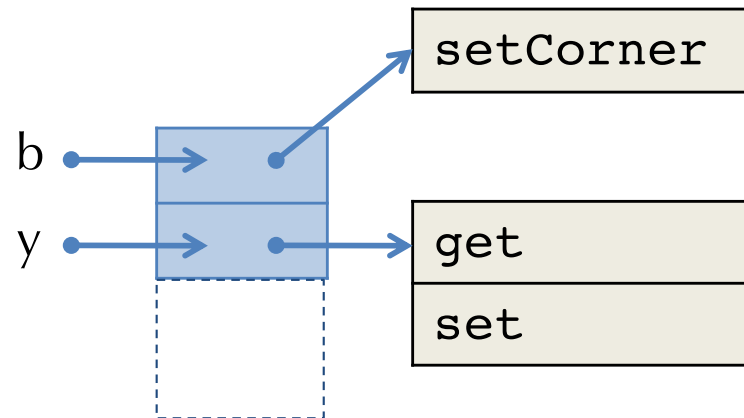
# Multiple Dispatch Vectors

- A reference to an object might have multiple “entry points”
  - Each entry point corresponds to a dispatch vector
  - Which one is used depends on the statically known type of the program.

```
Blob b = new Blob();  
Color y = b;    // implicit cast!
```

- Compile

```
Color y = b;  
As  
Movq [[b]] + 8 , y
```



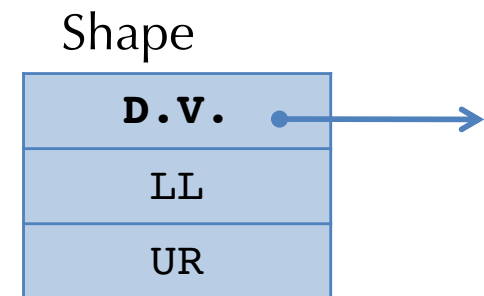
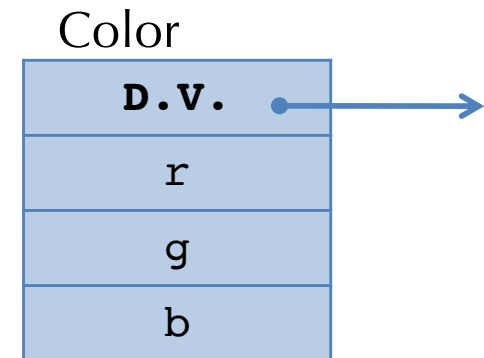
# Multiple D.V. Summary

- Benefit: Efficient dispatch, same cost as for multiple inheritance
- Drawbacks:
  - Cast has a runtime cost
  - More complicated programming model... hard to understand/debug?
- What about multiple inheritance and fields?

# Multiple Inheritance: Fields

- Multiple supertypes (Java): methods conflict (as we saw)
- Multiple inheritance (C++): fields can also conflict
- Location of the object's fields can no longer be a constant offset from the start of the object.

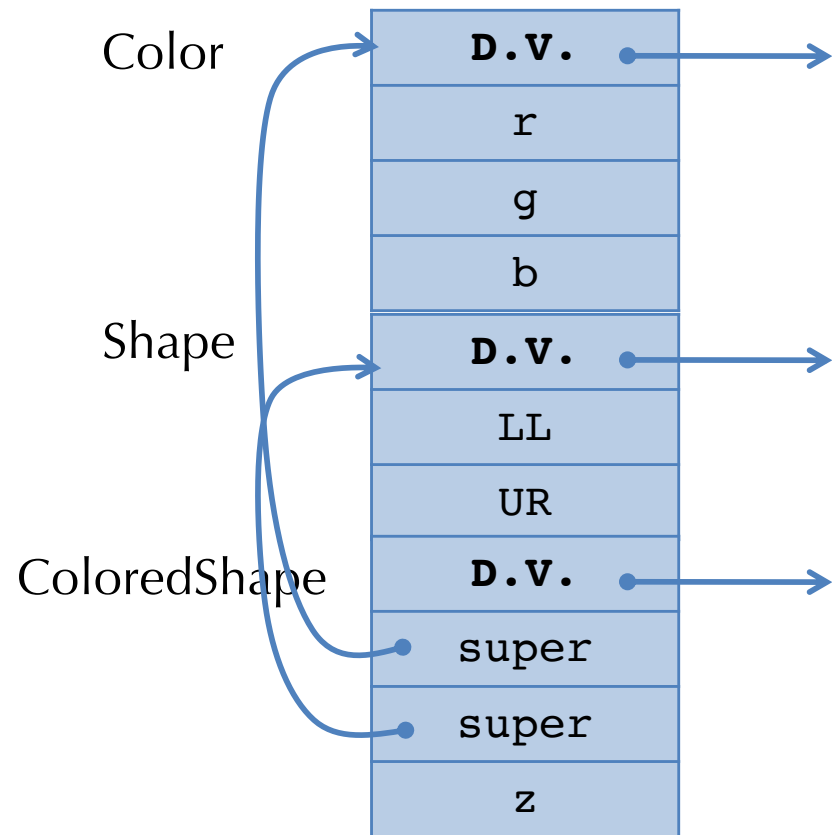
```
class Color {  
    float r, g, b; /* offsets: 4,8,12 */  
}  
class Shape {  
    Point LL, UR; /* offsets: 4, 8 */  
}  
class ColoredShape extends  
Color, Shape {  
    int z;  
}
```

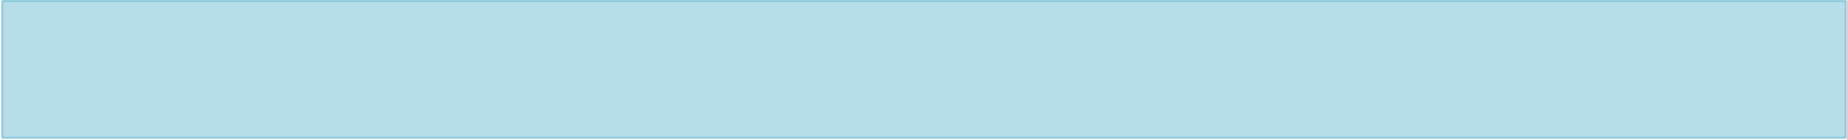


ColoredShape ??

# C++ approach:

- Add pointers to the superclass fields
  - Need to have multiple dispatch vectors anyway (to deal with methods)
- Extra indirection needed to access superclass fields
- Used even if there is a single superclass
  - Uniformity





Compiling lambda calculus to straight-line code.  
Representing evaluation environments at runtime.

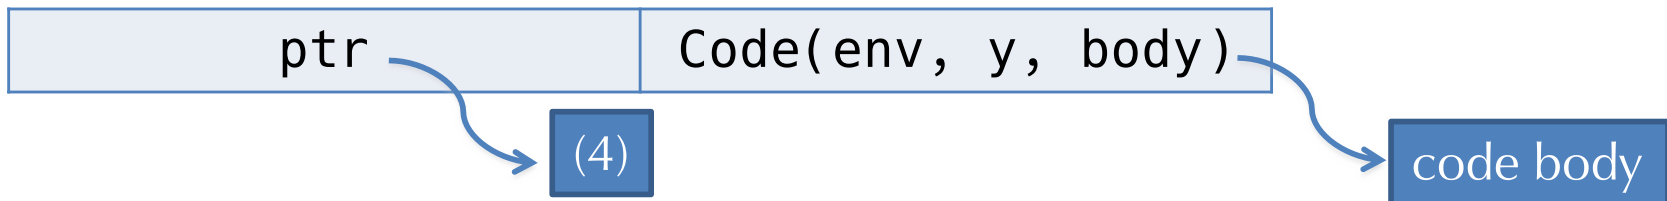
# CLOSURE CONVERSION REVISITED

# Compiling First-class Functions

- To implement first-class functions on a processor, there are two problems:
  - First: we must implement substitution of free variables
  - Second: we must separate 'code' from 'data'
- Reify the substitution:
  - Move substitution from the meta language to the object language by making the data structure & lookup operation explicit
  - The environment-based interpreter is one step in this direction
- Closure Conversion:
  - Eliminates free variables by packaging up the needed environment in the data structure.
- Hoisting:
  - Separates code from data, pulling closed code to the top level.

# Example of closure creation

- Recall the “add” function:  
`let add = fun x -> fun y -> x + y`
- Consider the inner function: `fun y -> x + y`
- When run the function application: `add 4`  
the program builds a closure and returns it.
  - The closure is a pair of the environment and a code pointer.



- The code pointer takes a pair of parameters: `env` and `y`
  - The function code is (essentially):  
`fun (env, y) -> let x = nth env 0 in x + y`

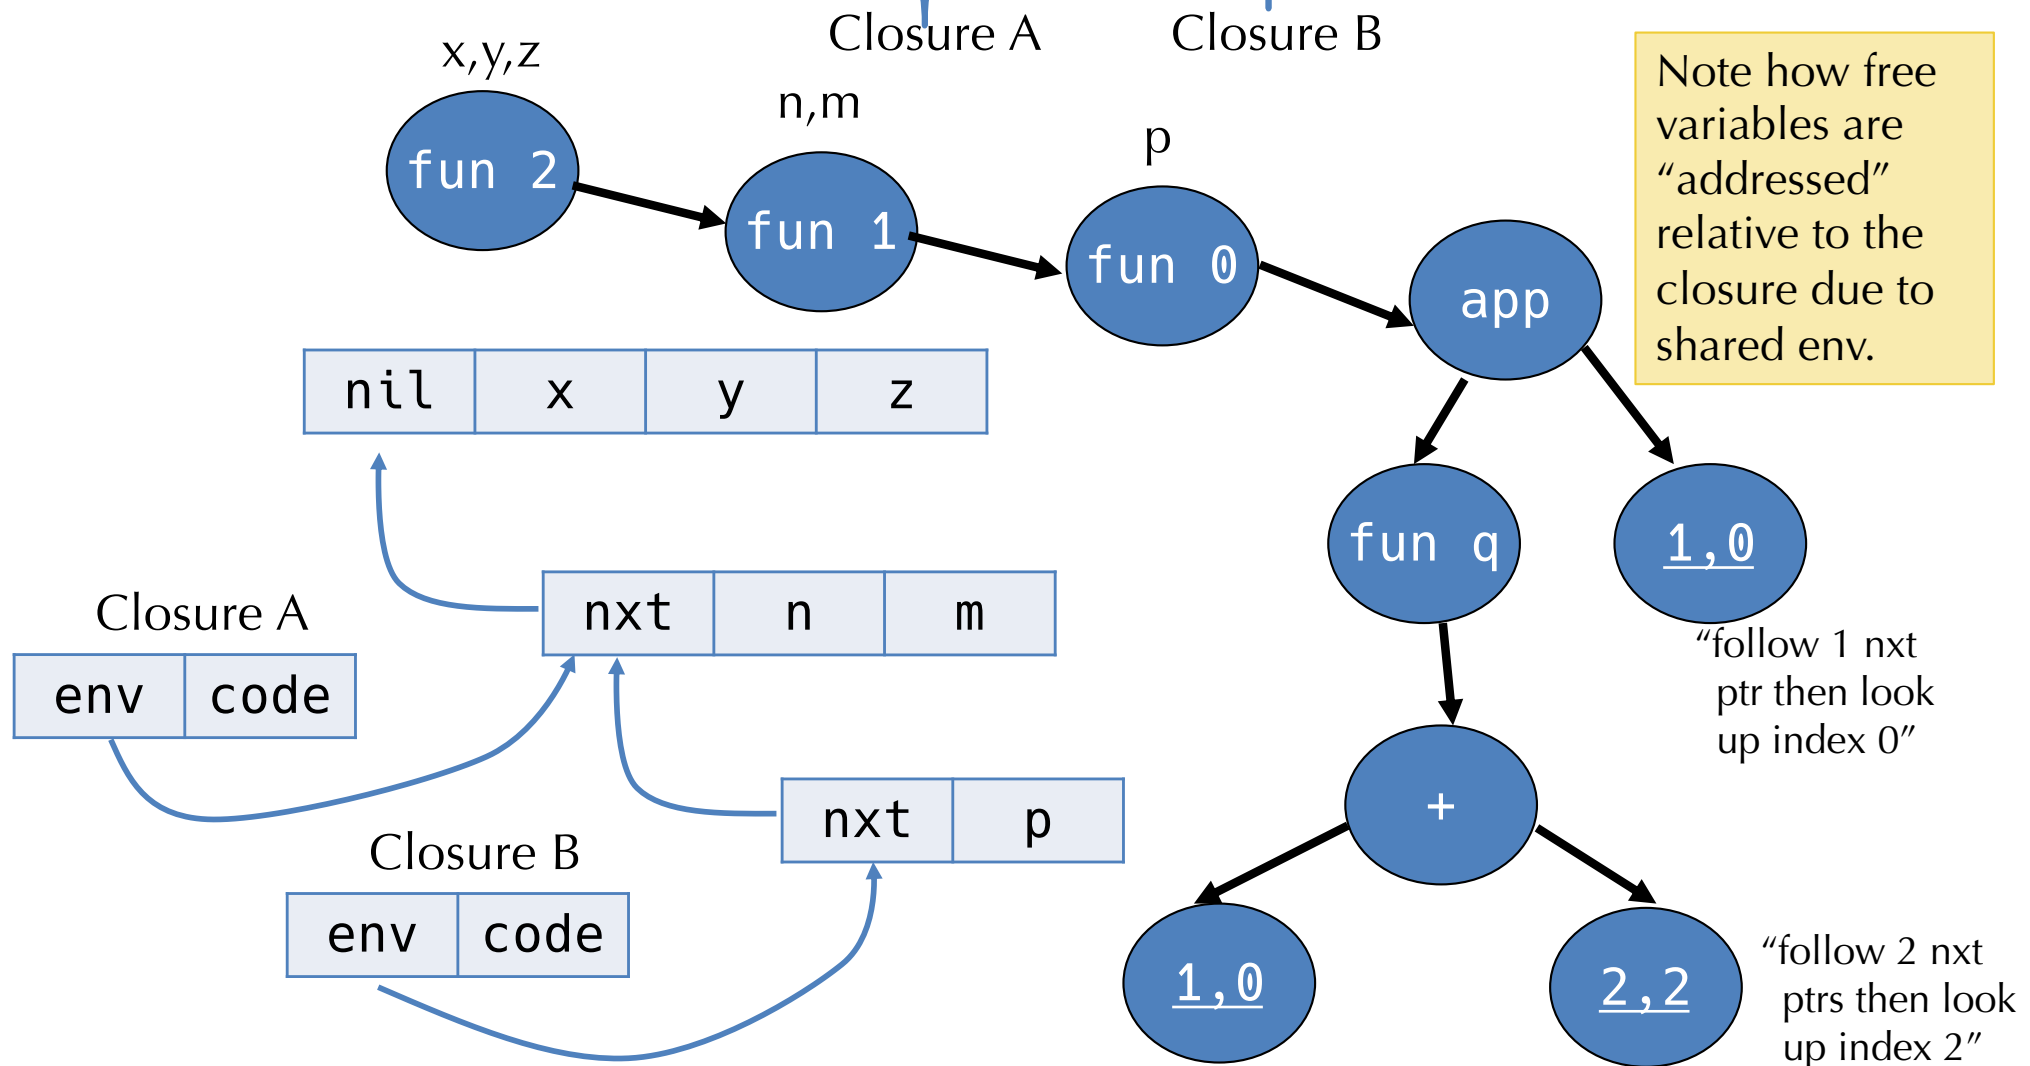


# Representing Closures

- As we saw, the simple closure conversion algorithm doesn't generate very efficient code.
  - It stores all the values for variables in the environment, even if they aren't needed by the function body.
  - It copies the environment values each time a nested closure is created.
  - It uses a linked-list datastructure for tuples.
- There are many options:
  - Store only the values for free variables in the body of the closure.
  - Share subcomponents of the environment to avoid copying
  - Use vectors or arrays rather than linked structures

# Array-based Closures with N-ary Functions

```
(fun (x y z) ->  
  (fun (n m) -> (fun p -> (fun q -> n + z) x)
```



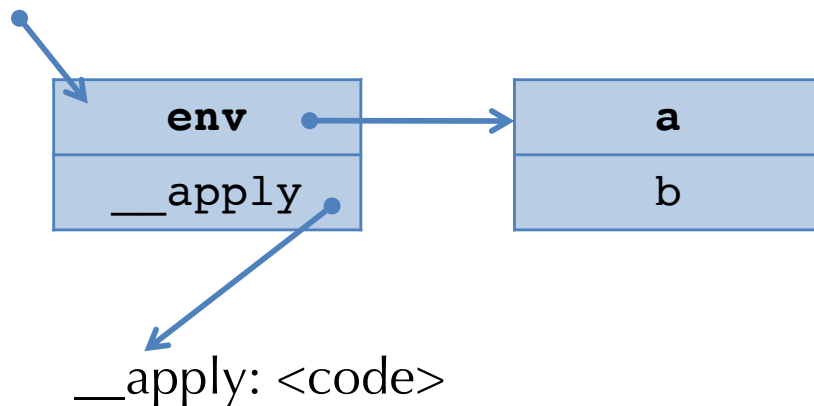
# Compiling Closures to LLVM IR

- The “types” of the environment data structures are generic tuples
  - The tuples contain a mix of int and closure values
  - We know statically what the tuple-type of the environment should be
  - LLVM IR doesn’t have generic types
- Type translations:
  - $\llbracket - \rrbracket$  for “interpretation” that retains type information
    - $\llbracket \text{int} \rrbracket = \text{i64}$
    - $\llbracket (t_1, \dots, t_n) \rrbracket = \{\llbracket t_1 \rrbracket, \dots, \llbracket t_n \rrbracket\}^*$
    - $\llbracket t_1 \rightarrow t_2 \rrbracket = \llbracket t_1 \rightarrow t_2 \rrbracket_C$
  - $\llbracket t_1 \rightarrow t_2 \rrbracket_C = \{\text{i8}^*, ((\text{i8}^*, \llbracket t_1 \rrbracket) \rightarrow \llbracket t_2 \rrbracket)^*\}^*$  “Closure Representation”
- Rough sketch:
  - Allocation & uses of objects use the “interpretation” translation
  - Anywhere an environment is passed or stored, use  $\text{i8}^*$  and bitcast to/from the translation type.

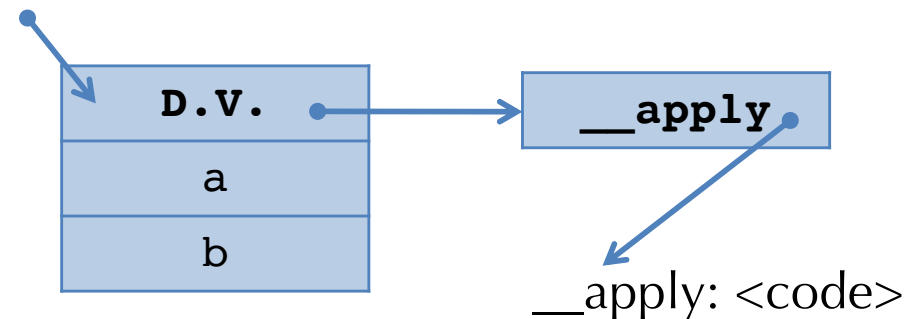
# Observe: Closure $\approx$ Single-method Object

- Free variables  $\approx$  Fields
- Environment pointer  $\approx$  “this” parameter
- Closure for function:  $\approx$  Instance of this class:

```
fun (x,y) ->  
  x + y + a + b
```



```
class C {  
  int a, b;  
  int apply(x,y) {  
    x + y + a + b  
  }  
}
```





Why optimize?

# OPTIMIZATIONS, GENERALLY

# Optimizations

- The code generated by our OAT compiler so far is pretty inefficient.
  - Lots of redundant moves.
  - Lots of unnecessary arithmetic instructions.
- Consider this OAT program:

```
int foo(int w) {  
    var x = 3 + 5;  
    var y = x * w;  
    var z = y - 0;  
    return z * 4;  
}
```

frontend.ml

```
define i64 @foo(i64 %w1) {  
    %w2 = alloca i64  
    %x5 = alloca i64  
    %y10 = alloca i64  
    %z14 = alloca i64  
    store i64 %w1, i64* %w2  
    %_bop4 = add i64 3, 5  
    store i64 %_bop4, i64* %x5  
    %x7 = load i64, i64* %x5  
    %w8 = load i64, i64* %w2  
    %_bop9 = mul i64 %x7, %w8  
    store i64 %_bop9, i64* %y10  
    %y12 = load i64, i64* %y10  
    %_bop13 = sub i64 %y12, 0  
    store i64 %_bop13, i64* %z14  
    %z16 = load i64, i64* %z14  
    %_bop17 = mul i64 %z16, 4  
    ret i64 %_bop17  
}
```

- opt-example.c, opt-example.oat

# Unoptimized vs. Optimized Output

```
define i64 @foo(i64 %_w1) {  
  %_w2 = alloca i64  
  %_x5 = alloca i64  
  %_y10 = alloca i64  
  %_z14 = alloca i64  
  store i64 %_w1, i64* %_w2  
  %_bop4 = add i64 3, 5  
  store i64 %_bop4, i64* %_x5  
  %_x7 = load i64, i64* %_x5  
  %_w8 = load i64, i64* %_w2  
  %_bop9 = mul i64 %_x7, %_w8  
  store i64 %_bop9, i64* %_y10  
  %_y12 = load i64, i64* %_y10  
  %_bop13 = sub i64 %_y12, 0  
  store i64 %_bop13, i64* %_z14  
  %_z16 = load i64, i64* %_z14  
  %_bop17 = mul i64 %_z16, 4  
  ret i64 %_bop17  
}
```

backend.ml

```
.text  
.globl _foo  
_foo:  
  pushq %rbp  
  movq %rsp, %rbp  
  subq $136, %rsp  
  movq %rdi, %rax  
  movq %rax, -8(%rbp)  
  pushq $0  
  movq %rsp, -16(%rbp)  
  pushq $0  
  movq %rsp, -24(%rbp)  
  pushq $0  
  movq %rsp, -32(%rbp)  
  pushq $0  
  movq %rsp, -40(%rbp)  
  movq -8(%rbp), %rcx  
  movq -16(%rbp), %rax  
  movq %rcx, (%rax)  
  movq $3, %rax  
  movq $5, %rcx  
  addq %rcx, %rax  
  movq %rax, -56(%rbp)  
  movq -56(%rbp), %rcx  
  movq -24(%rbp), %rax  
  movq %rcx, (%rax)  
  movq -24(%rbp), %rax  
  movq (%rax), %rcx  
  movq %rcx, -72(%rbp)  
  movq -16(%rbp), %rax  
  movq (%rax), %rcx  
  movq %rcx, -80(%rbp)  
  movq -72(%rbp), %rax  
  movq -80(%rbp), %rcx  
  imulq %rcx, %rax  
  movq %rax, -88(%rbp)  
  movq -88(%rbp), %rcx  
  movq -32(%rbp), %rax  
  movq %rcx, (%rax)  
  movq -32(%rbp), %rax  
  movq (%rax), %rcx  
  movq %rcx, -104(%rbp)  
  movq -104(%rbp), %rax  
  movq $0, %rcx  
  subq %rcx, %rax  
  movq %rax, -112(%rbp)  
  movq -112(%rbp), %rcx  
  movq -40(%rbp), %rax  
  movq %rcx, (%rax)  
  movq -40(%rbp), %rax  
  movq (%rax), %rcx  
  movq %rcx, -128(%rbp)  
  movq -128(%rbp), %rax  
  movq $4, %rcx  
  imulq %rcx, %rax  
  movq %rax, -136(%rbp)  
  movq -136(%rbp), %rax  
  movq %rbp, %rsp  
  popq %rbp  
  retq
```

Optimized code:

```
_foo:  
  pushq %rbp  
  movq %rsp, %rbp  
  movq %rdi, %rax  
  shlq $5, %rax  
  popq %rbp  
  retq
```

- Code above generated by `clang -O3`
- Function `foo` may be inlined by the compiler, so it can be implemented by just one instruction!

# Why do we need optimizations?

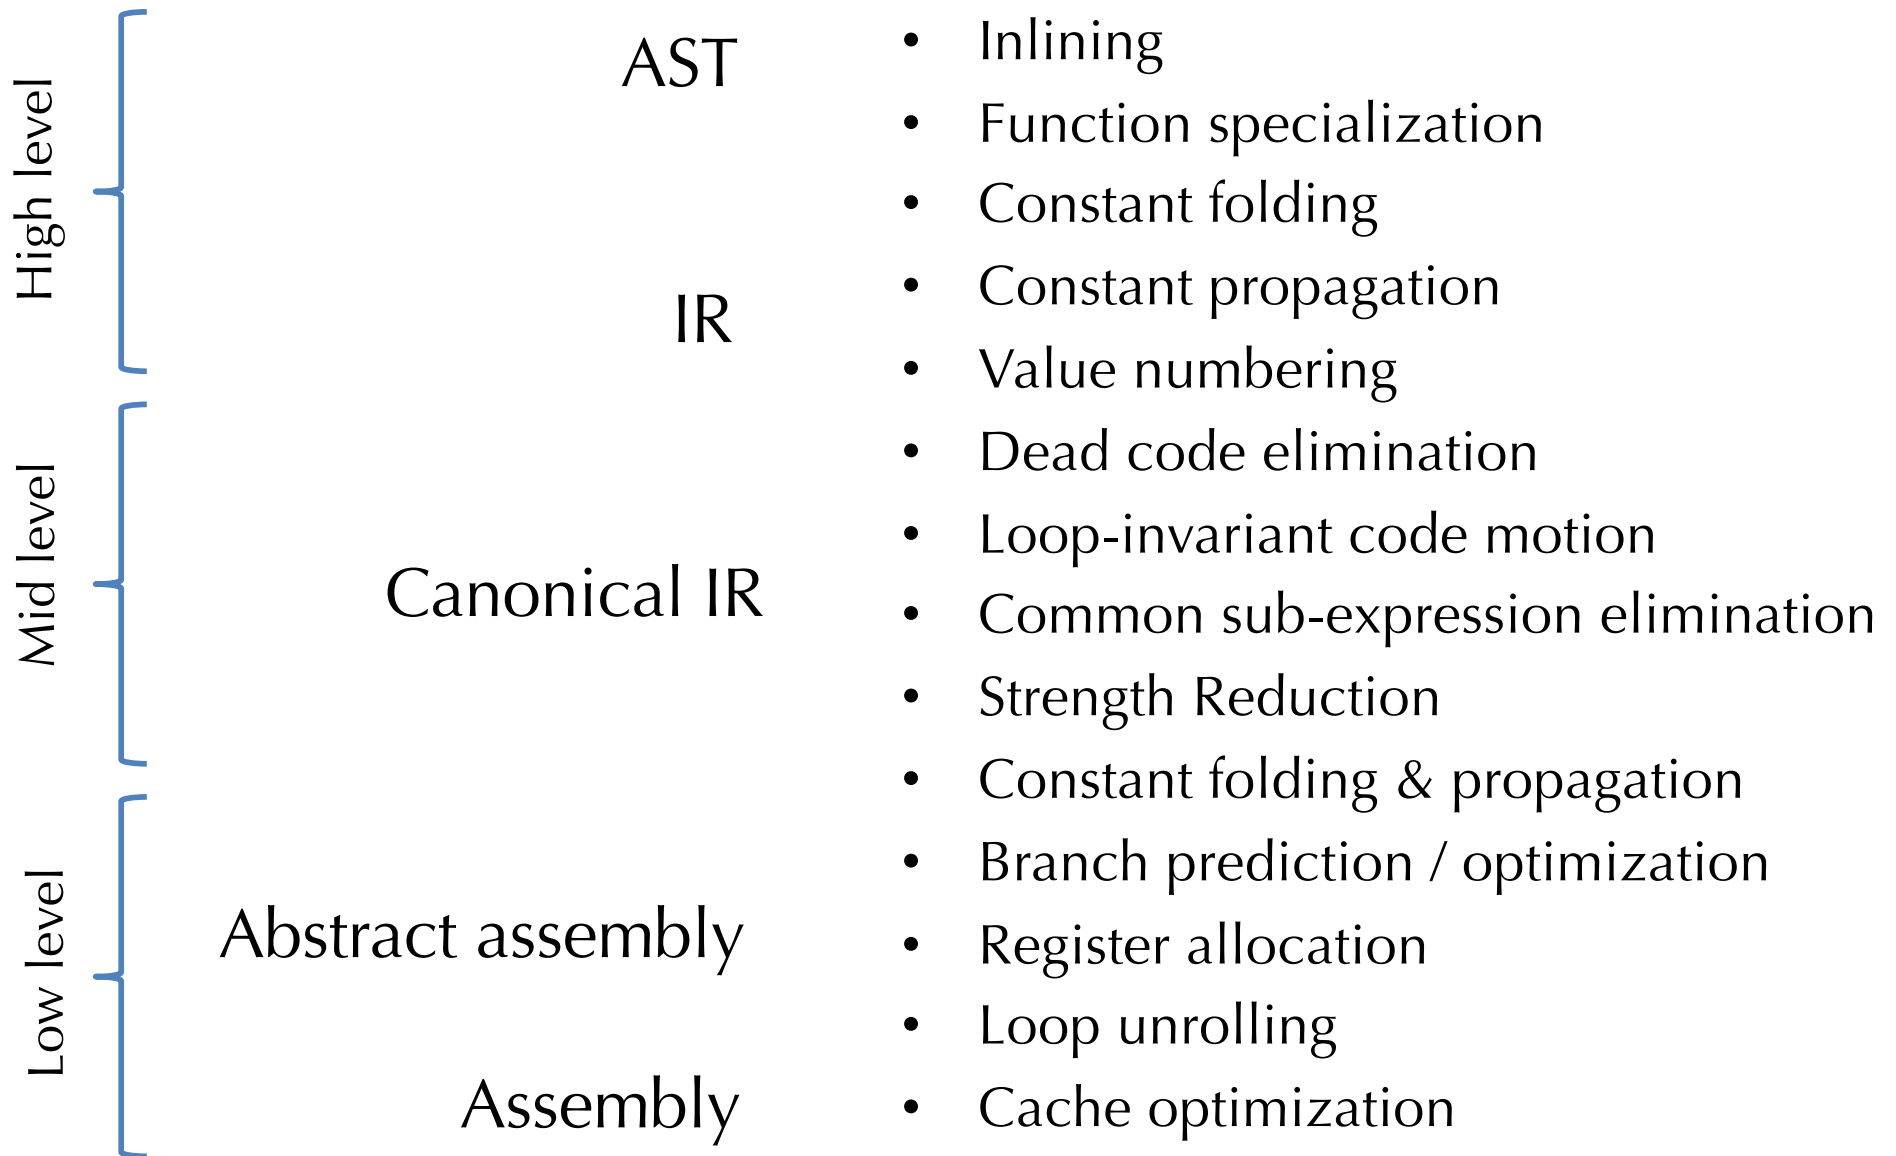
- To help programmers...
  - They write modular, clean, high-level programs
  - Compiler generates efficient, high-performance assembly
- Programmers don't write optimal code
- High-level languages make avoiding redundant computation inconvenient or impossible
  - e.g.  $A[i][j] = A[i][j] + 1$
- Architectural independence
  - Optimal code depends on features not expressed to the programmer
  - Modern architectures *assume* optimization
- Different kinds of optimizations:
  - Time: improve execution speed
  - Space: reduce amount of memory needed
  - Power: lower power consumption (e.g. to extend battery life)



# Some caveats

- Optimization are code transformations:
  - They can be applied at any stage of the compiler
  - They must be *safe* – they shouldn't change the meaning of the program.
- In general, optimizations require some program analysis:
  - To determine if the transformation really is safe
  - To determine whether the transformation is cost effective
- This course: most common and valuable performance optimizations
  - See Muchnick (optional text) for ~10 chapters about optimization

# When to apply optimization



# Where to Optimize?

- Usual goal: improve time performance
- Problem: many optimizations trade space for time
- Example: *Loop unrolling*
  - Idea: rewrite a loop like:  

```
for(int i=0; i<100; i=i+1) {  
    s = s + a[i];  
}
```
  - Into a loop like:  

```
for(int i=0; i<99; i=i+2){  
    s = s + a[i];  
    s = s + a[i+1];  
}
```
- Tradeoffs:
  - Increasing code space slows down whole program a tiny bit (extra instructions to manage) but speeds up the loop a lot
  - For frequently executed code with long loops: generally a win
  - Interacts with instruction cache and branch prediction hardware
- Complex optimizations may never pay off!

# Writing Fast Programs In Practice

- Pick the right algorithms and data structures.
  - These have a much bigger impact on performance than compiler optimizations.
  - Reduce # of operations
  - Reduce memory accesses
  - Minimize indirection – it breaks working-set coherence
- *Then* turn on compiler optimizations
- Profile to determine program hot spots
- Evaluate whether the algorithm/data structure design works
- ...if so: “tweak” the source code until the optimizer does “the right thing” to the machine code

# Safety

- Whether an optimization is *safe* depends on the programming language semantics.
  - Languages that provide weaker guarantees to the programmer permit more optimizations but have more ambiguity in their behavior.
  - e.g. In Java tail-call optimization (that turns recursive function calls into loops) is not valid.
  - e.g. In C, loading from initialized memory is undefined, so the compiler can do anything.
- Example: *loop-invariant code motion*
  - Idea: hoist invariant code out of a loop

```
while (b) {  
    z = y/x;  
    ...           // y, x not updated  
}
```



```
z = y/x;  
while (b) {  
    ...           // y, x not updated  
}
```

- Is this more efficient?
- Is this safe?



A high-level tour of a variety of optimizations.

# BASIC OPTIMIZATIONS

# Constant Folding

- Idea: If operands are known at compile time, perform the operation statically.

```
int x = (2 + 3) * y  ➔ int x = 5 * y  
b & false           ➔ false
```

- Performed at every stage of optimization...
- Why?
  - Constant expressions can be created by translation or earlier optimizations
- Example: `A[2]` might be compiled to:  
`MEM[MEM[A] + 2 * 4] ➔ MEM[MEM[A] + 8]`

# Constant Folding Conditionals

```
if (true) S           → S
if (false) S          → ;
if (true) S else S'   → S
if (false) S else S'  → S'
while (false) S       → ;

if (2 > 3) S          → ;
```



# Algebraic Simplification

- More general form of constant folding
  - Take advantage of mathematically sound simplification rules
- Identities:
  - $a * 1 \rightarrow a$                        $a * 0 \rightarrow 0$
  - $a + 0 \rightarrow a$                        $a - 0 \rightarrow a$
  - $b | \text{false} \rightarrow b$                $b \& \text{true} \rightarrow b$
- Reassociation & commutativity:
  - $(a + 1) + 2 \rightarrow a + (1 + 2) \rightarrow a + 3$
  - $(2 + a) + 4 \rightarrow (a + 2) + 4 \rightarrow a + (2 + 4) \rightarrow a + 6$
- Strength reduction: (replace expensive op with cheaper op)
  - $a * 4 \rightarrow a \ll 2$
  - $a * 7 \rightarrow (a \ll 3) - a$
  - $a / 32767 \rightarrow (a \gg 15) + (a \gg 30)$
- Note 1: must be careful with floating point (due to rounding) and integer arithmetic (due to overflow/underflow)
- Note 2: iteration of these optimizations is useful... how much?

# Constant Propagation

- If the value is known to be a constant, replace the use of the variable by the constant
- Value of the variable must be propagated forward from the point of assignment
  - This is a substitution operation

- Example:

```
int x = 5;  
int y = x * 2; → int y = 5 * 2; → int y = 10; →  
int z = a[y];    int z = a[y];    int z = a[y];  int z = a[10];
```

- To be most effective, constant propagation should be interleaved with constant folding

# Copy Propagation

- If one variable is assigned to another, replace uses of the assigned variable with the copied variable.
- Need to know where copies of the variable propagate.
- Interacts with the scoping rules of the language.

- Example:

```
x = y;  
if (x > 1) {  
    x = x * f(x - 1);  
}  
→  
x = y;  
if (y > 1) {  
    x = y * f(y - 1);  
}
```

- Can make the first assignment to *x* *dead* code (that can be eliminated).

# Dead Code Elimination

- If a side-effect free statement can never be observed, it is safe to eliminate the statement.

```
x  = y * y  // x is dead!  
...          // x never used  →      ...  
x  = z * z          x  = z * z
```

- A variable is *dead* if it is never used after it is defined.
  - Computing such *definition* and *use* information is an important component of compiler
- Dead variables can be created by other optimizations...

# Unreachable/Dead Code

- Basic blocks not reachable by any trace leading from the starting basic block are *unreachable* and can be deleted.
  - Performed at the IR or assembly level
  - Improves cache, TLB performance
- Dead code: similar to unreachable blocks.
  - A value might be computed but never subsequently used.
- Code for computing the value can be dropped
- But only if it's *pure*, i.e. it has *no externally visible side effects*
  - Externally visible effects: raising an exception, modifying a global variable, going into an infinite loop, printing to standard output, sending a network packet, launching a rocket
  - Note: Pure functional languages (e.g. Haskell) make reasoning about the safety of optimizations (and code transformations in general) easier!

# Inlining

- Replace a call to a function with the body of the function itself with arguments rewritten to be local variables:
- Example in OAT code:

```
int g(int x) { return x + pow(x); }  
int pow(int a) { int b = 1; int n = 0;  
    while (n < a) {b = 2 * b}; return b; }
```



```
int g(int x) { int a = x; int b = 1; int n = 0;  
    while (n < a) {b = 2 * b}; tmp = b; return x + tmp;  
}
```

- May need to rename variable names to avoid *name capture*
  - Example of what can go wrong?
- Best done at the AST or relatively high-level IR.
- When is it profitable?
  - Eliminates the stack manipulation, jump, etc.
  - Can increase code size.
  - Enables further optimizations

# Code Specialization

- Idea: create specialized versions of a function that is called from different places with different arguments.

- Example: specialize function `f` in:

```
class A implements I { int m() {...} }  
class B implements I { int m() {...} }  
int f(I x) { x.m(); }           // don't know which m  
A a = new A(); f(a);           // know it's A.m  
B b = new B(); f(b);           // know it's B.m
```

- `f_A` would have code specialized to dispatch to `A.m`
- `f_B` would have code specialized to dispatch to `B.m`
- You can also inline methods when the run-time type is known statically
  - Often just one class implements a method.

# Common Subexpression Elimination

- In some sense it's the opposite of inlining: fold redundant computations together
- Example:

$a[i] = a[i] + 1$  compiles to:

$[a + i*4] = [a + i*4] + 1$

Common subexpression elimination removes the redundant add and multiply:

$t = a + i*4; [t] = [t] + 1$

- For safety, you must be sure that the shared expression always has the same value in both places!



# Unsafe Common Subexpression Elimination

- Example: consider this OAT function:

```
unit f(int[] a, int[] b, int[] c) {  
    int j = ...; int i = ...; int k = ...;  
    b[j] = a[i] + 1;  
    c[k] = a[i];  
    return;  
}
```

- The optimization that shares the expression `a[i]` is unsafe... why?

```
unit f(int[] a, int[] b, int[] c) {  
    int j = ...; int i = ...; int k = ...;  
    t = a[i];  
    b[j] = t + 1;  
    c[k] = t;  
    return;  
}
```