

# CS 516: COMPILERS

## Lecture 8

### *Topics*

- Data structures in LLVM
- LLVMlite and Homework 4

### *Materials*

- `lec08.zip`



# DATATYPES IN THE LLVM IR

# Structured Data in LLVM

- LLVM's IR uses types to describe the structure of data.

```
t ::=  
  void  
  i1 | i8 | i64           N-bit integers  
  [<#elts> x t]          arrays  
  fty                     function types  
  {t1, t2, ... , tn}    structures  
  t*                      pointers  
  %Tident                 named (identified) type  
  
fty ::=                   Function Types  
  t (t1, ..., tn)       return, argument types
```

- <#elts> is an integer constant  $\geq 0$
- Structure types can be named at the top level:

```
%T1 = type {t1, t2, ... , tn}
```

- Such structure types can be recursive

# Example LL Types

- An array of 516 integers: `[ 516 x i64 ]`
- A two-dimensional array of integers: `[ 3 x [ 4 x i64 ] ]`
- Structure for representing arrays with their length:  
`{ i64 , [ 0 x i64 ] }`
  - There is no array-bounds check; the static type information is only used for calculating pointer offsets.
- C-style linked lists (declared at the top level):  
`%Node = type { i64, %Node* }`
- Structs from the C program shown earlier:  
`%Rect = { %Point, %Point, %Point, %Point }`  
`%Point = { i64, i64 }`

# IR for types

**Source:**

```
struct point { int x; int y; }  
struct rect { point ll; point lr; point ul; point ur }  
...  
myR.ul.x = 42;  
int t = myR.ur.y;
```




**IR:**

```
%Rect = { %Point, %Point, %Point, %Point }  
%Point = { i64, i64 }  
...  
(something) = 42; // store a new ul.x  
...  
%13 = (something); // load the ur y
```



**X86:**

```
...  
Mov $42, (X86 address)  
Mov (X86 address), %eax
```



# getelementptr

- LLVM provides the `getelementptr` instruction to compute pointer values
  - Given a pointer and a “path” through the structured data pointed to by that pointer, `getelementptr` computes an address
  - `getelementptr` **does not access memory**.
  - This is the abstract analog of the X86 LEA (load effective address).
  - It is a “type indexed” operation, since the size computations depend on the type

```
insn ::= ...  
      | getelementptr t* %val, t1 idx1, t2 idx2 ,...
```

- Example: access the x component of the first point of a rectangle:

```
%tmp1 = getelementptr %Rect* %square, i32 0, i32 0  
%tmp2 = getelementptr %Point* %tmp1, i32 0, i32 0
```

# GEP Example\*

```
struct RT {  
    int A;  
    int B[10][20];  
    int C;  
}  
struct ST {  
    struct RT X;  
    int Y;  
    struct RT Z;  
}  
int *foo(struct ST *s) {  
    return &s[1].Z.B[5][13];  
}
```



```
%RT = type { i32, [10 x [20 x i32]], i32 }  
%ST = type { %RT, i32, %RT }  
define i32* @foo(%ST* %s) {  
entry:  
    %arrayidx = getelementptr %ST* %s, i32 1, i32 2, i32 1, i32 5, i32 13  
    ret i32* %arrayidx  
}
```

# GEP Example\*

```
struct RT {  
    int A;  
    int B[10][20];  
    int C;  
}  
struct ST {  
    struct RT X;  
    int Y;  
    struct RT Z;  
}  
int *foo(struct ST *s) {  
    return &s[1].Z.B[5][13];  
}
```

1. %s is a pointer to an (array of) %ST structs,  
suppose the pointer value is ADDR

```
%RT = type { i32, [10 x [20 x i32]], i32 }  
%ST = type { %RT, i32, %RT }  
define i32* @foo(%ST* %s) {  
entry:  
    %arrayidx = getelementptr %ST* %s, i32 1, i32 2, i32 1, i32 5, i32 13  
    ret i32* %arrayidx  
}
```



# GEP Example\*

```
struct RT {  
    int A;  
    int B[10][20];  
    int C;  
}  
struct ST {  
    struct RT X;  
    int Y;  
    struct RT Z;  
}  
int *foo(struct ST *s) {  
    return &s[1].Z.B[5][13];  
}
```

1. %s is a pointer to an (array of) %ST structs, suppose the pointer value is ADDR

2. Compute the index of the 1<sup>st</sup> element by adding size\_ty(%ST).

```
%RT = type { i32, [10 x [20 x i32]], i32 }  
%ST = type { %RT, i32, %RT }  
define i32* @foo(%ST* %s) {  
entry:  
    %arrayidx = getelementptr %ST* %s, i32 1, i32 2, i32 1, i32 5, i32 13  
    ret i32* %arrayidx  
}
```

# GEP Example\*

```
struct RT {  
    int A;  
    int B[10][20];  
    int C;  
}  
struct ST {  
    struct RT X;  
    int Y;  
    struct RT Z;  
}  
int *foo(struct ST *s) {  
    return &s[1].Z.B[5][13];  
}
```

1. %s is a pointer to an (array of) %ST structs, suppose the pointer value is ADDR

2. Compute the index of the 1<sup>st</sup> element by adding `size_ty(%ST)`.

3. Compute the index of the Z field by adding `size_ty(%RT) + size_ty(i32)` to skip past X and Y.

```
%RT = type { i32, [10 x [20 x i32]], i32 }  
%ST = type { %RT, i32, %RT }  
define i32* @foo(%ST* %s) {  
entry:  
    %arrayidx = getelementptr %ST* %s, i32 1, i32 2, i32 1, i32 5, i32 13  
    ret i32* %arrayidx  
}
```

# GEP Example\*

```
struct RT {
    int A;
    int B[10][20];
    int C;
}
struct ST {
    struct RT X;
    int Y;
    struct RT Z;
}
int *foo(struct ST *s) {
    return &s[1].Z.B[5][13];
}
```

1. %s is a pointer to an (array of) %ST structs, suppose the pointer value is ADDR

2. Compute the index of the 1<sup>st</sup> element by adding `size_ty(%ST)`.

3. Compute the index of the Z field by adding `size_ty(%RT) + size_ty(i32)` to skip past X and Y.

4. Compute the index of the B field by adding `size_ty(i32)` to skip past A.

```
%RT = type { i32, [10 x [20 x i32]], i32 }
%ST = type { %RT, i32, %RT }
define i32* @foo(%ST* %s) {
entry:
    %arrayidx = getelementptr %ST* %s, i32 1, i32 2, i32 1, i32 5, i32 13
    ret i32* %arrayidx
}
```

# GEP Example\*

```
struct RT {
    int A;
    int B[10][20];
    int C;
}
struct ST {
    struct RT X;
    int Y;
    struct RT Z;
}
int *foo(struct ST *s) {
    return &s[1].Z.B[5][13];
}
```

1. %s is a pointer to an (array of) %ST structs, suppose the pointer value is ADDR

2. Compute the index of the 1<sup>st</sup> element by adding `size_ty(%ST)`.

3. Compute the index of the Z field by adding `size_ty(%RT) + size_ty(i32)` to skip past X and Y.

4. Compute the index of the B field by adding `size_ty(i32)` to skip past A.

5. Index into the 2d array.

```
%RT = type { i32, [10 x [20 x i32]], i32 }
%ST = type { %RT, i32, %RT }
define i32* @foo(%ST* %s) {
entry:
    %arrayidx = getelementptr %ST* %s, i32 1, i32 2, i32 1, i32 5, i32 13
    ret i32* %arrayidx
}
```

# GEP Example\*

```
struct RT {
    int A;
    int B[10][20];
    int C;
}
struct ST {
    struct RT X;
    int Y;
    struct RT Z;
}
int *foo(struct ST *s) {
    return &s[1].Z.B[5][13];
}
```

1. %s is a pointer to an (array of) %ST structs, suppose the pointer value is ADDR

2. Compute the index of the 1<sup>st</sup> element by adding `size_ty(%ST)`.

3. Compute the index of the Z field by adding `size_ty(%RT) + size_ty(i32)` to skip past X and Y.

4. Compute the index of the B field by adding `size_ty(i32)` to skip past A.

5. Index into the 2d array.

```
%RT = type { i32, [10 x [20 x i32]], i32 }
%ST = type { %RT, i32, %RT }
define i32* @foo(%ST* %s) {
entry:
    %arrayidx = getelementptr %ST* %s, i32 1, i32 2, i32 1, i32 5, i32 13
    ret i32* %arrayidx
}
```

Final answer:  $\text{ADDR} + \text{size\_ty}(\%ST) + \text{size\_ty}(\%RT) + \text{size\_ty}(i32) + \text{size\_ty}(i32) + 5 \cdot 20 \cdot \text{size\_ty}(i32) + 13 \cdot \text{size\_ty}(i32)$

# getelementptr

- GEP *never* dereferences the address it's calculating:
  - GEP only produces pointers by doing arithmetic
  - It doesn't actually traverse the links of a datastructure
- To index into a deeply nested structure, need to “follow the pointer” by loading from the computed pointer
  - See list.ll from HW4

```

%node = type { i64, %node* }

@hd = global %node { i64 1, %node* @md }
@md = global %node { i64 2, %node* @tl }
@tl = global %node { i64 3, %node* null }

define i64 @main(i64 %argc, i8** %argv) {
    %head = getelementptr %node, %node* @hd, i32 0, i32 0
    %link = getelementptr %node, %node* @hd, i32 0, i32 1
    %next = load %node*, %node** %link
    %val = getelementptr %node, %node* %next, i32 0, i32 0
    %link2 = getelementptr %node, %node* %next, i32 0, i32 1
    %next2 = load %node*, %node** %link2
    %val2 = getelementptr %node, %node* %next2, i32 0, i32 0
    %1 = load i64, i64* %val2
    ret i64 %1
}

```

hw4/llprograms/list1.ll

# Compiling Datastructures via LLVM

1. Translate high level language *types* into an *LLVM type*.
  - For some languages (e.g. C) this process is straight forward
    - The translation simply uses platform-specific alignment and padding
  - For other languages, (e.g. OO languages) there might be a fairly complex elaboration.
    - e.g. for Ocaml, arrays types might be translated to pointers to length-indexed structs.

`[[int array]] = { i32, [0 x i32]}*`

2. Translate *accesses* of the data into *getelementptr* operations:

- e.g. for Ocaml array size access:

`[[length a]] =`

`%1 = getelementptr {i32, [0xi32]}* %a, i32 0, i32 0`





see HW4: `lib/ll/ll.ml`

# LLVMLITE SPECIFICATION

# LLVMlite notes

- Real LLVM requires that constants appearing in `getelementptr` be declared with type `i32`:

```
%struct = type { i64, [5 x i64], i64}

@gbl = global %struct {i64 1,
    [5 x i64] [i64 2, i64 3, i64 4, i64 5, i64 6], i64 7}

define void @foo() {
    %1 = getelementptr %struct* @gbl, i32 0, i32 0
    ...
}
```

- LLVMlite ignores the `i32` annotation and treats these as `i64` values
  - we keep the `i32` annotation in the syntax to retain compatibility with the clang compiler



# COMPILING LLVM-LITE TO X86

# Compiling LLVM locals

# Compiling LLVM locals

- How do we manage storage for each %uid defined by an LLVM instruction?

# Compiling LLVM locals

- How do we manage storage for each %uid defined by an LLVM instruction?
- Option 1:
  - Map each %uid to a x86 register
  - Efficient!
  - Difficult to do effectively: many %uid values, only 16 registers
  - We will see how to do this later in the semester

# Compiling LLVM locals

- How do we manage storage for each %uid defined by an LLVM instruction?
- Option 1:
  - Map each %uid to a x86 register
  - Efficient!
  - Difficult to do effectively: many %uid values, only 16 registers
  - We will see how to do this later in the semester
- Option 2:
  - Map each %uid to a stack-allocated space
  - Less efficient!
  - Simple to implement

# Compiling LLVM locals

- How do we manage storage for each %uid defined by an LLVM instruction?
- Option 1:
  - Map each %uid to a x86 register
  - Efficient!
  - Difficult to do effectively: many %uid values, only 16 registers
  - We will see how to do this later in the semester
- Option 2:
  - Map each %uid to a stack-allocated space
  - Less efficient!
  - Simple to implement
- For HW3 we will follow Option 2



# Compiling LLVMlite Types to X86

- `[[i1]], [[i64]], [[t*]]` = quad word (8 bytes, 8-byte aligned)
- raw `i8` values are not allowed (they must be manipulated via `i8*`)
- array and struct types are laid out sequentially in memory
- `getelementptr` computations must be relative to the LLVMlite size definitions
  - i.e. `[[i1]]` = quad

# Other LLVMlite Features

- Calls
  - Follow x64 AMD ABI calling conventions
  - Should interoperate with C programs
- More types: structured data records and arrays
- New instruction: `getelementptr`
  - LLVM IR's way of dealing with structured data
  - trickiest part of the compilation process
  - Note: you can start HW3 before understanding `getelementptr`
- New instruction: `bitcast`
  - convert between pointer types
- Globals
  - must use `%rip` relative addressing. See next slide...

# RIP-relative global variables

```
Mov global_foo(%rip)
```

# 2<sup>64</sup> total bytes

0x00000000

01101100

1 byte

0x00000008

```
0xffffffff
```

1 quadword

8 bytes

---

64 bits

# Announcements

- HW4: LLVM lite
  - **Goal:** “Backend” compiler from LLVMlite → X86lite
  - **Available:** Later today on the course web pages.
  - **Due:** Thursday, March 2nd at 11:59:59pm
  - **Teams:** Teams of 2. Only one group member needs to submit

***START EARLY!!***



see HW4 and README

ll.ml, using oatc, clang, etc.

## TOUR OF HW 4

# Bitcast

- What if the LLVM IR's type system isn't expressive enough?
  - e.g. if the source language has subtyping, perhaps due to inheritance
  - e.g. if the source language has polymorphic/generic types
- LLVM IR provides a `bitcast` instruction
  - This is a form of (potentially) unsafe cast. Misuse can cause serious bugs (segmentation faults, or silent memory corruption)

```
%rect2 = type { i64, i64 }           ; two-field record
%rect3 = type { i64, i64, i64 }      ; three-field record

define @foo() {
    %1 = alloca %rect3               ; allocate a three-field record
    %2 = bitcast %rect3* %1 to %rect2* ; safe cast
    %3 = getelementptr %rect2* %2, i32 0, i32 1 ; allowed
    ...
}
```